

Rosetta

A User's Guide

Multi-language bindings for unmodified C++,
generated once with C++26 static reflection

Frantz Maerten
`fmaerten@gmail.com`

September 3, 2026

Preface

This book is about a single, narrow claim: that the information needed to expose a C++ library to Python, to JavaScript, to Lua, to a REST endpoint or to a property panel is *already in the header*, and that a compiler can be made to read it out.

Everything else follows from taking that claim seriously. You do not write a binding. You write down *what* you want bound — a list of headers and a list of target languages — and a program compiled against your types answers the rest by reflection. Your class definitions never change.

Who this is for

You have a C++ library. You want it callable from somewhere else, ideally from several somewheres, and you would rather not maintain a hand-written wrapper per destination. You are comfortable with CMake and with your own build.

You do *not* need to know anything about C++26 static reflection to use rosetta. Reflection is what makes the tool possible, and Chapter 2 explains where it runs and why the pipeline has the shape it does — but the surface you work with is a JSON file and a command-line tool.

How to read it

In a hurry

Chapter 3 is a complete first binding, from an unmodified header to an importable Python module. Read that, then come back.

Evaluating rosetta

Chapters 1 and 2 say what the system does and what it costs, including the parts that do not work.

Using it in earnest

Part II is the manifest, field by field, and is the part you will return to. Part IV is the tool that consumes it.

Pushing past bindings

Part V covers the runtime object model, scripting the metadata itself, and writing a backend of your own.

The appendices are lookup tables: every manifest field, every annotation, every target, and the errors you are most likely to hit.

Conventions

`Monospace` is a literal — a field name, a path, a command. A *manifest* always means the `manifest.json` driving one binding project. *Generation host* is the machine that runs the reflection compiler; *target* is wherever the generated binding is finally built, which is deliberately not the same machine.

Two callouts appear throughout:

NOTE

Something worth knowing that is not on the shortest path.

WATCH OUT

Something that will cost you an afternoon if you meet it unprepared.

Status

Rosetta is a prototype, and this book says so where it matters. It tracks in-flight C++ proposals — P2996 (reflection), P3394 (annotations), P3294 (token injection) — and is developed against the Bloomberg `clang-p2996` fork, which is currently the only front-end that carries all of what it uses. What the tool *generates* has no such constraint: that is the whole point, and Chapter 2 is where the distinction is made precise.

Related documents

The repository carries the reference material this book is built from — `docs/MANIFEST.md`, `docs/ROSETTA_GEN.md`, `docs/ANNOTATIONS.md`, `docs/COVERAGE.md` and the rest. Where a detail here is compressed, those are the long form. The academic write-up of the same system, with its evaluation and its related work, lives under `paper/`; it argues that reflection can serve as an interface description language, where this book simply assumes it and gets on with the job.

Contents

Preface	i
I Understanding rosetta	1
1 What rosetta is	2
1.1 The problem, stated once	2
1.2 The idea	2
1.3 What you write	3
1.4 What comes out without being asked for	3
1.5 What you get out	4
1.6 Reflection-free by construction	4
1.7 The nineteen targets	5
1.8 Annotations are optional	5
1.9 What rosetta is not	5
1.10 Where to go next	6
2 The pipeline, and why it has that shape	7
2.1 The constraint everything follows from	7
2.2 The five steps	7
2.3 Your library is read twice, for different reasons	8
2.4 Two commands, not five	9
2.5 When to re-run what	9
2.6 The footgun the tool warns about	10
3 Your first binding	11
3.1 Prerequisites	11
3.2 Build the tool, once	11
3.3 The library — leave it alone	11
3.4 The manifest	12
3.5 Generate and build — the one-command version	13
3.6 The same thing, one step at a time	13
3.7 Build and use it	14
3.8 Read <code>coverage.json</code> once	14
3.9 Enriching, if you want to	14
3.10 Cleaning up	15
3.11 What to read next	15

II	The manifest	16
4	Anatomy of a manifest	17
4.1	The minimum	17
4.2	Paths resolve from the manifest	17
4.3	Comments	18
4.4	Variables	18
4.5	Where the toolchain lives	19
4.6	Naming things	20
4.7	Shared defaults: <code>namespace</code> and <code>header_dir</code>	20
4.8	Multiple include directories	21
4.9	A field map	22
5	Declaring what to bind	23
5.1	Classes	23
5.2	Binding a whole folder	24
5.3	Renaming: <code>expose</code>	26
5.4	Free functions	26
5.5	Extension methods	27
5.6	Out-parameters	28
5.7	Module init	29
6	Targets	30
6.1	The two entry forms	30
6.2	The available <code>lang</code> values	30
6.3	Per-target linker flags	31
6.4	Where the built artifact lands	32
6.5	Pinning the runtime	32
6.6	Choosing targets	33
7	Making the generated projects build	34
7.1	The two ways to supply definitions	34
7.2	Preprocessor definitions	36
7.3	Build type and optimisation	36
7.4	Headers the library's own build produces	37
7.5	Python wheels	37
7.6	A realistic manifest	39
8	Types across the boundary	40
8.1	Foreign sequence containers	40
8.2	Foreign matrices	41
8.3	Foreign-library interop	42
8.4	Two standard types that need no declaration	43
8.5	Deciding what you need	44
9	Presets	45
9.1	Using one	45
9.2	What applying one does	46
9.3	The shipped preset	46
9.4	Writing your own	46

III	Annotations	48
10	Annotations	49
10.1	Two spellings, one result	49
10.2	The annotation set	50
10.3	Who consumes what	51
10.4	Out-of-line annotations	51
10.5	Internal annotations	53
10.6	Gotchas	53
IV	Driving the tool	54
11	The <code>rosetta_gen</code> tool	55
11.1	Building it	55
11.2	The four modes	55
11.3	<code>-init</code> : starting a project	56
11.4	<code>-build</code> : the whole pipeline	56
11.5	Plain mode: one step at a time	59
11.6	<code>-clean</code> : back to sources	59
11.7	Exit status	60
12	Overloads and coverage	61
12.1	Overloads	61
12.2	<code>coverage.json</code>	62
12.3	Using it	63
12.4	Instrumenting a backend of your own	63
12.5	A note for backend authors	64
V	Beyond bindings	65
13	The dynamic object model	66
13.1	What it emits	66
13.2	Using it	66
13.3	What the tables look like	66
13.4	Why <code>register_all()</code> exists	67
13.5	Calling by name	67
13.6	Two consequences worth knowing	68
13.7	What it buys	68
13.8	What it costs	68
13.9	When to reach for it	69
14	Scripting the metadata itself	70
14.1	The trick, in one sentence	70
14.2	Why a facade rather than binding <code>dynamic.h</code>	70
14.3	The one hand-written piece	71
14.4	What a project has to write	71
14.5	What it buys	72
14.6	Running the pipeline twice	72
15	Writing a backend	73

15.1	The mental model: two layers	73
15.2	Layer 1: the runtime visitor	73
15.3	Layer 2: the generation backend	75
15.4	Registering it	76
15.5	Wiring it into a project	76
15.6	API reference	77
15.7	Where to copy from	78
VI	Worked examples	79
16	Worked examples	80
16.1	The map	80
16.2	The smallest case	81
16.3	A geometry library	81
16.4	The dynamic model, end to end	81
16.5	Rosetta in the wild	82
16.6	Extending a generated binding by hand	82
VII	Reference	84
A	Manifest field reference	85
A.1	Class entry fields	87
A.2	Function entry fields	87
A.3	Target entry fields	87
A.4	<code>user_lib</code> entry fields	87
B	Annotation reference	88
B.1	The kinds	88
B.2	Widget hints	89
B.3	Consumption by backend family	89
B.4	Inline versus side-car	89
B.5	Side-car schema	89
B.6	Merge rules	90
B.7	Internal	90
C	Target reference	91
C.1	Every <code>lang</code> value	91
C.2	Feature support	91
C.3	Packaging	92
C.4	Interop and foreign types	92
C.5	Build shapes for <code>-build</code>	93
C.6	External requirements	93
D	Troubleshooting	94
D.1	A method is missing from the binding	94
D.2	The generated driver does not compile	94
D.3	A side-car annotation had no effect	95
D.4	The binding is slow	95
D.5	The module will not load	95
D.6	The pipeline produced stale output	95

D.7	-build skipped a backend	96
D.8	The wrong Python was used	96
D.9	-clean did not remove something	96
D.10	A variable in the manifest did not expand	96
D.11	Reporting something	96

Part I

Understanding rosetta

Chapter 1

What rosetta is

1.1 The problem, stated once

Exposing a C++ library to another language is normally paid for once per destination. A pybind11 module for Python. An N-API layer for Node. A sol2 usertype for Lua. A P/Invoke shim for C#. A moc-annotated hierarchy for a Qt panel. A JSON route table for a REST service.

Look at what those files actually contain. Here is one method, five times:

```
double Grid::at(int i, int j) const;
```

```
py::class_<Grid>(m, "Grid").def("at", &Grid::at);           // pybind11
ut["at"] = &Grid::at;                                       // sol2
exports.Set("at", Napi::Function::New(env, wrap_at));       // N-API
emscripten::class_<Grid>("Grid").function("at", &Grid::at); // embind
[DllImport("grid")] static extern double Grid_at(IntPtr h, int i, int j);
```


Every one of those restates, in a different notation, something the C++ compiler already knows: that `Grid` has a member called `at`, that it takes two ints, that it returns a `double`, that it does not mutate. The information is not missing. It is simply not reachable from anywhere except the compiler.

And the cost is not the initial write — it is the *derivative*. Add a parameter to `at` and five files are now wrong, in five languages, and four of them will not tell you until run time.

1.2 The idea

C++26 static reflection (P2996) makes that information reachable from inside C++ itself. A program can ask, at its own compile time, what members a type has, what their types are, whether they are public, whether they are virtual.

Rosetta does exactly one thing with that: it reads your headers *once*, through reflection, into a target-neutral intermediate representation, and then *projects* that representation into as many bindings as you asked for.



You declare *what* to bind. You never declare *how*.

1.3 What you write

Two things, and one of them you already have.

The header stays exactly as it is:

```
// my_lib/person.h
#include <string>

struct Person {
    std::string name;
    int         age = 0;
    std::string id;
    std::string greet(const std::string &salutation) const;
};
```

The manifest is the only new file:

```
{
  "user_include": "../my_lib",
  "rosetta_include": "/path/to/rosetta/include",
  "targets": [
    { "lang": "python", "name": "person_py" },
    { "lang": "node",   "name": "person_js" }
  ],
  "classes": [
    { "name": "Person", "header": "person.h" }
  ]
}
```

Then one command:

```
rosetta_gen --build manifest.json
```

and `bindings/python` and `bindings/node` exist, compiled, with `Person` in both — its three fields as properties, its method as a method, its constructors as constructors.

1.4 What comes out without being asked for

The manifest above names one class. Reflection supplies the rest of it:

- **Public fields**, as read/write properties with per-backend getters and setters.
- **Public methods**, instance and static alike.
- **Constructors** — every overload, not just the default one.

- **Inheritance.** Public base fields and methods are flattened into the derived binding; a derived declaration shadows the base one, and a virtual diamond collapses to a single member. Virtual methods are flagged, so backends that can generate trampolines do.
- **Overload sets.** The whole set reaches the IR; what each target does with it depends on whether its runtime can dispatch on argument types (Chapter 12).
- **Enums,** both plain and scoped, with their enumerators.
- **Nested user types and `std::vector`,** marshalled across the boundary — a `Surface` returning `std::vector<Triangle>` works without a line about `Triangle` beyond listing it.

Free functions are the one thing you must name, because a namespace has no member list to enumerate — you list them in the manifest, still without touching the header (Section 5.4).

1.5 What you get out

One self-contained project per target, under `bindings/`:

```
bindings/
python/   auto_pybind.cpp      CMakeLists.txt  pyproject.toml  make_wheel.py
node/     auto_napi.cpp        CMakeLists.txt  package.json
wasm/     auto_emscripten.cpp  CMakeLists.txt
typescript/ person.d.ts
coverage.json
```

Each is an ordinary CMake (or npm, or emcmake) project. Nothing in it refers to rosetta at run time; nothing in it needs reflection to compile. That last point is worth its own section.

1.6 Reflection-free by construction

Every compiled backend *fully expands* what it found. There are no splices left in the output, no `<experimental/meta>`, no trace that reflection was ever involved — just explicit `pybind11` / `N-API` / `embind` / `sol2` / `jlccxx` calls, or, for the `dynamic` target, aggregate-initialised tables.

The consequence is the one that makes rosetta usable in practice:

NOTE

Reflection runs *once*, on the generation host. The machine that builds the generated binding needs an off-the-shelf toolchain — plain Clang, GCC, MSVC, a stock emsdk, a stock Qt 6. It does not need a C++26 compiler, and it does not need rosetta.

So the workflow for a project whose CI runs on a normal compiler is: generate on one machine that has the fork, commit or ship the generated tree, and build it anywhere.

Two targets are exceptions and are marked as such throughout: `rest` and `json` emit code that still splices reflections, so those two — and only those two — need the C++26 toolchain at their own build time.

1.7 The nineteen targets

Kind	lang	Output
Scripting	python	pybind11 extension module
	nanobind	nanobind extension module
	node	N-API native addon
	wasm	Emscripten / embind module
	lua	sol2 module, require -able
	julia	CxxWrap.jl / jlccxx shared module
Managed	csharp	C-ABI shared library + P/Invoke wrappers
	java	C-ABI + handle-backed FFM wrappers
User inter-	qt	Qt Widgets property/method inspector
faces	qml	QML / QtQuick inspector
	imgui	Dear ImGui inspector application
Services	rest	cpp-http lib JSON server + browser client
	openapi	OpenAPI 3.1 specification
	json	nlohmann (de)serialisation
Documents	typescript	ambient <code>.d.ts</code> declarations
	markdown	API reference document
	html	styled API reference page
	paraview	ParaView Server Manager plugin XML
Metadata	dynamic	the IR as runtime data (Chapter 13)

The full table, with what each one needs to build, is Appendix C. The last row is not a language binding at all and is the subject of its own chapter.

1.8 Annotations are optional

Reflection binds the class. *Annotations* enrich it — a docstring, a numeric range enforced at the boundary, a read-only property, a UI hint that turns a field into a slider:

```
[[= rosetta::doc{"Max solver iterations"}, = rosetta::range{0, 200},
  = rosetta::widget::slider ]]
int maxIter{100};
```

They are never required, they never change the C++ behaviour of your type, and you add them only where they pay off. If you would rather not touch the header at all — because it is a third party's, or because you want the generated binding to build on a stock compiler — the same metadata goes in a JSON side-car instead, wired by one manifest field. Chapter 10 covers both forms.

1.9 What rosetta is not

Being clear about this early saves disappointment later.

- **It is not a C++ parser.** It does not read your headers as text. A compiler does, and

rosetta asks the compiler. The one place a heuristic scan exists — `rosetta_gen -init` over a source tree — is explicitly a starting point you then edit.

- **It is not a runtime.** Nothing rosetta ships is loaded by your finished binding. The generated code depends on pybind11, or on N-API, or on sol2 — not on rosetta.
- **It does not bind everything.** Some C++ has no shape in the destination language: a `std::function` parameter, a raw pointer return, an operator with no name to bind to. Those members are *skipped*, never fatal, and every skip is recorded with a reason in `coverage.json` (Chapter 12). A binding that silently loses a method is the failure mode this design exists to avoid.
- **It is not finished.** It is a prototype tracking proposals that are themselves not final. See the status note in the Preface.

1.10 Where to go next

Chapter 2 explains why generating a binding is a *pipeline* rather than a command — which is the one piece of rosetta’s shape that is not obvious, and the piece that explains most of the rest. If you would rather see it work first, skip to Chapter 3.

Chapter 2

The pipeline, and why it has that shape

Binding a library with *rosetta* is a *pipeline* rather than a single command. That is not an implementation choice — it is the only shape the problem admits. This chapter explains why, because almost every surprise in later chapters resolves back to it.


2.1 The constraint everything follows from

C++26 reflection answers questions **at the compile time of a program that includes your headers**. It cannot be queried from a script, from a configuration file, or from a build system. Only from C++, being compiled, with your types in scope.

So to learn the shape of your types, something has to **compile**. To turn those answers into files on disk, that something has to **run**.

Compile, then run. Everything below is that sequence, plus one step before it and one after.

2.2 The five steps



Step 0 — build `rosetta_gen`

`rosetta_gen` is a JSON-to-C++ text templater. It does no reflection at all and builds with an off-the-shelf C++17 compiler. You build it once.

It exists so that what *you* write is a manifest and not a driver program. Given `manifest.json` it emits three files: a `<name>.cpp` containing a `main()` that reflects, a `bindings.h` with your types (and any out-of-line annotations) baked in, and a `CMakeLists.txt` that knows where the reflection toolchain lives.

Your library: untouched. `rosetta_gen` never opens your headers. It only copies their names into the driver it writes.

Step 1 — compile and run the driver

Three commands, because “compile then run” is three things — emit the program, build it, execute it:

```
rosetta_gen manifest.json gen      # emit the driver source
cmake -S gen -B gen/build          # COMPILE it -- reflection happens HERE
cmake --build gen/build
./generator bindings               # RUN it -- writes one project per target
```

During the *compile*, clang-p2996 evaluates the reflection operators in the driver: it walks your classes, enumerates fields and methods, reads annotations, and resolves each member against the target language’s type gate. **Every decision about what can and cannot be bound is made at this moment.** During the *run*, the driver does nothing clever — it prints the conclusions it was compiled with.

Neither half can be skipped. A compiler cannot write your binding files; a program cannot see types it was not compiled against.

Your library: read for the **first** time, as source, by the C++26 toolchain.

Step 2 — what lands in bindings/

One self-contained project per target, plus `coverage.json` — a machine-readable account of what bound and what did not, available *before* anything is compiled.

The generated code is fully expanded (Section 1.6). No splices, no `<experimental/meta>`, nothing that needs the fork.

Step 3 — compile a backend

Ordinary CMake, or npm, or emcmake, over an ordinary C++ project.

Your library: read a **second** time — the generated binding `#includes` your headers and calls your functions. This is where `user_sources` are compiled and `user_lib` is linked in.

Step 4 — load it

`import`, `require`, `dlopen`. Nothing rosetta-specific is left.

2.3 Your library is read twice, for different reasons

Pass	By	To	Needs C++26
step 1	the reflection driver	describe your types	yes
step 3	the generated binding	call them	no

This table is the single most useful thing to hold in your head, and it has one consequence worth stating on its own:

| WATCH OUT

Annotations written **inline** — `[[= rosetta::doc{...} ]]` — pull `<experimental/meta>` into your header. Pass 2 then needs the C++26 toolchain as well, and the “no reflection on the target” promise is lost. Out-of-line annotations (a `.ann.json` side-car) keep your header plain C++ and pass 2 reflection-free. See Section 10.4.

2.4 Two commands, not five

You rarely run the steps by hand. `rosetta_gen -build` chains all of them:

```
rosetta_gen --build manifest.json
```

It emits `gen/`, builds and runs the driver, and then compiles every backend the manifest declares — each with its own build shape (plain `cmake`, `npm install + npm run build`, `emcmake`, a second-stage `dotnet build` or `mvn package`). A backend whose toolchain is missing is *skipped with a note*, not an error, and every run ends with a per-backend summary.

`rosetta_gen -clean manifest.json` takes you back to sources. Chapter 11 is the full reference for both.

The manual sequence remains useful for exactly the reasons you would expect: inspecting `bindings.h`, hacking on the emitted driver, or wiring the stages into a build system of your own.

2.5 When to re-run what

The two-pass structure decides this, and getting it wrong is the most common way to be confused by rosetta.

What changed	What to do
A field or method: added, removed, retyped	Rebuild the backend. The generated binding is expanded, so this needs a regeneration <i>if the member set changed</i> — re-run <code>-build</code> .
An inline annotation	Re-run <code>-build</code> : the annotation is read during pass 1.
A side-car <code>.ann.json</code>	Re-run <code>rosetta_gen</code> <i>and</i> the generator, not just the binding build — the JSON is baked into <code>bindings.h</code> at <code>rosetta_gen</code> time.
A class added, or targets changed	Re-run the whole pipeline.
A <code>user_sources</code> glob now matches new files	Re-run <code>rosetta_gen</code> : globs expand at generation time.
The manifest moved	Re-run. Every relative path in it resolves from the manifest’s own directory.

| NOTE

Re-emitting is idempotent: unchanged files are not rewritten, so the next `cmake -build` recompiles nothing. Running the pipeline again when you are not sure is cheap.

2.6 The footgun the tool warns about

If the emitted sources changed but a `generator` binary from a previous emit still sits next to the output directory, running it without rebuilding **silently re-emits the old bindings**. The tool detects this and prints a rebuild note. Believe it.

The related one: a fresh checkout (or a `-clean`) leaves `out_dir/build` unconfigured, so `cmake -build` alone fails. The *configure* step is printed as a `next:` hint when that happens.

Chapter 3

Your first binding

This chapter goes from an unmodified header to an importable Python module and a loadable Node addon. It is deliberately complete: every command, in order, with what each one produces.

3.1 Prerequisites

- A **clang-p2996** build. Rosetta looks for it at `$ENV{HOME}/devs/c++/clang-p2996/build` by default; anywhere else is fine, you just say where (Section 4.5).
- **CMake** 3.28 or later, and **Ninja** or **Make**.
- For the Python target, `pip install pybind11`. For Node, `npm`. For WebAssembly, an activated `emsdk`.

The reflection flags rosetta’s generated projects use, should you need to know them: `-freflection -freflection-latest -fexperimental-library`, plus `-fannotation-attributes` for anything using annotations.

| NOTE

Only the *generation host* needs the fork. Everything under `bindings/` builds with a stock compiler — with the two documented exceptions, `rest` and `json`.

3.2 Build the tool, once

```
cmake -G Ninja -S tools/rosetta_gen -B tools/rosetta_gen/build
cmake --build tools/rosetta_gen/build
```

The binary always lands in `<repo>/bin/rosetta_gen` — not in the build tree. This step needs no reflection: `rosetta_gen` is a JSON reader and a text templater, and builds with an off-the-shelf C++17 compiler.

3.3 The library — leave it alone

```
// my_lib/person.h
#pragma once
#include <string>
```

```
struct Person {
    std::string name;
    int         age = 0;
    std::string id;

    Person() = default;
    Person(std::string n, int a) : name(std::move(n)), age(a) {}

    std::string greet(const std::string &salutation) const;
};
```

Note what is *not* there: no macro, no registration call, no rosetta include, no export annotation. That is the whole point.

3.4 The manifest

Put it wherever you like — every relative path inside resolves from the manifest's own directory. Here, `my_lib/rosetta/manifest.json`:

```
{
  "user_include": "..",
  "rosetta_include": "/path/to/rosetta/include",
  "generator_name": "person_gen",
  "module_name": "person",

  "targets": [
    { "lang": "python", "name": "person_py" },
    { "lang": "node",   "name": "person_js" },
    "typescript",
    "markdown"
  ],

  "classes": [
    { "name": "Person", "header": "person.h" }
  ]
}
```

Four required ideas and nothing else:

`user_include`

where your headers live. One directory, or an array of them.

`rosetta_include`

where rosetta's `include/` lives.

`targets`

which backends to emit. A bare string uses `module_name`; the object form sets a per-target module name.

`classes`

what to bind. `header` is required; `name` defaults to the header's stem.

`generator_name` is optional — it names the emitted driver and defaults to the manifest's folder name.

NOTE

Do not want to write this by hand? `rosetta_gen -init` writes a fully-commented starter manifest, and `rosetta_gen -init manifest.json ./src` pre-fills it from a scan of your source tree. See Section 11.3.

3.5 Generate and build — the one-command version

```
/path/to/rosetta/bin/rosetta_gen --build manifest.json
```

That emits the generator project, builds it (this is where reflection runs), runs it, and then compiles every backend. The run ends with a summary:

```
== summary
python      OK
node        OK
typescript  OK (nothing to compile)
markdown    OK (nothing to compile)
```

A backend whose toolchain is missing is skipped with a note rather than failing the run:

```
wasm        SKIPPED (emcc not found -- activate emsdk)
```

3.6 The same thing, one step at a time

Worth doing once, because it makes the pipeline concrete.

```
# 1. manifest -> generator project (no reflection yet)
rosetta_gen manifest.json gen

# 2. build the generator -- REFLECTION HAPPENS HERE
cmake -S gen -B gen/build && cmake --build gen/build

# 3. run it -- one self-contained project per target
./generator bindings
```

After step 1, look inside `gen/`. There are three files:

<code>person_gen.cpp</code>	a <code>main()</code> that calls <code>rosetta::generate<Person>(...)</code> — the driver
<code>bindings.h</code>	your headers included, plus any baked-in out-of-line annotations
<code>CMakeLists.txt</code>	knows where clang-p2996 is and which flags it needs

After step 3:

```
bindings/
  coverage.json
  python/      auto_pybind.cpp  CMakeLists.txt  pyproject.toml  make_wheel.py
  node/        auto_napi.cpp    CMakeLists.txt  package.json
  typescript/  person.d.ts
  markdown/    person.md
```

Open `auto_pybind.cpp`. It is ordinary pybind11 — explicit `.def()` calls, one per member, with no reflection anywhere. That file is shippable: it compiles on any machine with pybind11 and a normal C++ compiler.

3.7 Build and use it

Python

```
cd bindings/python
cmake -G Ninja -B build && cmake --build build
```

```
import person_py

p = person_py.Person("Alice", 30)
p.name = "Alice B."
print(p.greet("Hello"))
print(p.age)
```

Node

```
cd bindings/node
npm install && npx cmake-js compile
```

```
const m = require('./build/Release/person_js.node');
const p = new m.Person('Alice', 30);
console.log(p.greet('Hello'));
```

The two you did not have to build

`bindings/typescript/person.d.ts` declares the Node addon's surface, so an editor autocompletes it. `bindings/markdown/person.md` is an API reference page. Neither compiles anything — they are text, produced directly.

3.8 Read `coverage.json` once

Before moving on, open `bindings/coverage.json`. Even for a three-field class it tells you something worth having: exactly which members reached each backend and which did not, with a reason for every drop.

```
{
  "rosetta_coverage": 1,
  "reflection": [],
  "targets": [
    { "target": "python", "bound": 5, "skipped": 0, "classes": [ ... ] }
  ]
}
```

On a real library it will not be empty, and the habit of checking it is the difference between knowing your binding surface and guessing at it. Chapter 12 is about reading it properly.

3.9 Enriching, if you want to

Everything above used reflection alone: names match the C++ identifiers, types map automatically, and you get no docstrings, no validation and no read-only semantics. Adding them is opt-in, member by member:

```
#include <rosetta/annotations.h>

struct Person {
    [[= rosetta::doc{"display name"}]]
    std::string name;

    [[= rosetta::doc{"age in years"}, = rosetta::range{0.0, 150.0}]]
    int age = 0;

    [[= rosetta::doc{"server-assigned id"}, = rosetta::readonly{}]]
    std::string id;
};
```

Now `help(person_py.Person.name)` says something, assigning 200 to `age` raises, and `id` has no setter.

| WATCH OUT

Writing annotations inline makes the header C++26-only — every translation unit that includes it, including the generated binding, now needs the fork. If you want to keep the header plain C++ (and you usually do), use the JSON side-car instead: same metadata, same result, one manifest field. See Section 10.4.

3.10 Cleaning up

```
rosetta_gen --clean manifest.json
```

removes `gen/`, the generator binary, and the per-backend folders under `bindings/` *that this manifest declares*. Anything else you parked there survives, and a folder whose `CMakeLists.txt` lacks the `Generated by rosetta` stamp is never touched. The manifest itself is never removed.

3.11 What to read next

You now have a working binding. Part II is the manifest in full — and the next chapter starts with the rules that apply to all of it.

Part II

The manifest

Chapter 4

Anatomy of a manifest

The manifest is the single hand-written file that drives rosetta. This chapter covers the rules that apply to *all* of it — how paths resolve, how to write comments, how to factor out repetition — before the following chapters go field by field.

It answers three questions for the framework:

1. **Where** are your headers, and where is rosetta's `include/`?
2. **What** classes and free functions should be bound?
3. **Which** language backends should be emitted?

Everything else — fields, methods, constructors, enums, inheritance — is discovered by reflection from the headers themselves.

4.1 The minimum

Three keys are required: `user_include`, `rosetta_include` and `targets`, plus `classes` unless a preset supplies it.

```
{
  "user_include": "./geom",
  "rosetta_include": ".././include",
  "targets": ["python"],
  "classes": [ { "header": "Point.h" } ]
}
```

That is a complete, working manifest.

4.2 Paths resolve from the manifest

Every relative path in the file is resolved against **the directory holding the manifest**, not the current working directory and not the output tree. Absolute paths are taken as written.

| WATCH OUT

Move the manifest and you must re-run `rosetta_gen`. Nothing detects a stale resolution for you.

There is one deliberate exception, and it is the one people trip over: a class entry's `header` is resolved against `user_include`, not against the manifest. That is because the header is emitted into an `#include "..."` line in generated code — it must be spelled relative to the *include path*, exactly as the compiler will look for it. `user_sources`, which name files to compile rather than to include, do resolve from the manifest.

4.3 Comments

JSON has no comment syntax, so the manifest uses two conventions, and both work everywhere:

```
{
  "/*": "Bindings for the geom library.",
  "/*paths": "All relative paths resolve from THIS file's directory.",

  // Line comments are tolerated by the parser too.
  /* As are block comments. */
  "user_include": "./geom"
}
```

Any key beginning with `//` is ignored — `"/1"`, `"/note"`, `"/cpp26"`. Keys must still be unique, which is why the numbered and named forms exist. The `//` and `/* */` line forms are accepted by the parser as an extension; the key form is plain JSON and survives any tool that reads the file.

4.4 Variables

A path that appears in several fields would otherwise be written out in each of them. The `cpp26_*` block is the case that motivated this — one toolchain root, spelled four times with a different suffix:

```
"cpp26_root": "$ENV{HOME}/devs/c++/clang-p2996/build",
"cpp26_cxx":  "$ENV{HOME}/devs/c++/clang-p2996/build/bin/clang++",
"cpp26_cc":   "$ENV{HOME}/devs/c++/clang-p2996/build/bin/clang",
"cpp26_lib":  "$ENV{HOME}/devs/c++/clang-p2996/build/lib",
```

`variables` declares a name for the shared part, which `$NAME` (or `${NAME}`) then stands for:

```
"variables": [
  { "name": "P2996", "value": "$ENV{HOME}/devs/c++/clang-p2996" },
],

"cpp26_root": "$P2996/build",
"cpp26_cxx":  "$P2996/build/bin/clang++",
"cpp26_cc":   "$P2996/build/bin/clang",
"cpp26_lib":  "$P2996/build/lib",
```

The substitution runs over **every string in the manifest** before any field is read, so a variable works anywhere a string does — `user_include`, `out_dir`, a `user_lib` directory, a `compile_definitions` value — not only in the toolchain paths above. The field itself disappears once applied.

Declarations are an **array**, not an object, because they are **ordered**: a value may use the variables declared before it, and only those.

```
"variables": [
  { "name": "TC", "value": "/opt/toolchains/clang-p2996" },
```

```
{ "name": "BIN", "value": "$TC/build/bin" }
],
"cpp26_cxx": "$BIN/clang++"
```

4.4.1 What is not substituted

A manifest already carries two other dollar notations, and both belong to **CMake**, which expands them at its own configure time. Rosetta hands them over untouched, so the rule is that **only a name you declared is substituted**:

Written	Result
<code>\$ENV{HOME}</code>	passed through verbatim — CMake's environment lookup, whatever <code>variables</code> says
<code>\${CMAKE_SOURCE_DIR}</code>	passed through verbatim — no <code>CMAKE_SOURCE_DIR</code> was declared
<code>\$P2996</code> with <code>P2996</code> declared	substituted

Declaring `ENV` is an error, since `$ENV{...}` could never resolve to it.

WATCH OUT

The price of that leniency is that a **misspelt variable is silent**. `$P2996x` is one name, `P2996x`, which nothing declares — so it stays as written and surfaces later as a path that does not exist. `${P2996}x` is how a variable abuts other text, and is worth preferring for that reason.

A variable name is a C identifier: letters, digits and `_`, not starting with a digit.

4.5 Where the toolchain lives

Four optional fields point at the C++26 / P2996 reflection compiler used to build the generator. Omit them all and rosetta uses these defaults:

Field	Default
<code>cpp26_root</code>	<code>\$ENV{HOME}/devs/c++/clang-p2996/build</code>
<code>cpp26_cxx</code>	<code>\${cpp26_root}/bin/clang++</code>
<code>cpp26_cc</code>	<code>\${cpp26_root}/bin/clang</code>
<code>cpp26_lib</code>	<code>\${cpp26_root}/lib</code>

If your clang-p2996 build lives elsewhere, set `cpp26_root` alone — the other three derive from it. Override an individual one only when the compiler binaries or the `libc++` / `libc++abi` directory sit outside the usual `bin/` and `lib/` layout.

`$ENV{HOME}` is expanded by CMake at configure time, so the path stays portable across machines. Of the targets, these fields affect only `rest` and `json` — every other one builds with an off-the-shelf compiler and ignores them entirely. (The *generator* always needs them.)

The command line wins over the manifest: `rosetta_gen -build -clang-p2996-root PATH`.

`qt_dir` plays the same role for the `qt` and `qml` backends, with `-qt-dir` as its flag.

4.6 Naming things

`generator_name`

Base name of the emitted reflection driver: `"my_person_gen"` produces `gen/my_person_gen.cpp`. The built binary is always called `generator`, so this is mostly cosmetic — its one load-bearing role is being the last-resort default for `module_name`. Defaults to the manifest's directory name.

`module_name`

The default binding module name, used by any bare-string target. A preset may supply one, which is why a preset-based manifest often names neither this nor `generator_name`.

NOTE

The *derived* `generator_name` is folded to an identifier (`my-cool-lib` becomes `my_cool_lib`) so a directory name cannot produce an invalid module name. A value you spell out is used verbatim — if you write something that is not an identifier, you get it.

4.7 Shared defaults: `namespace` and `header_dir`

When every entry repeats the same namespace and the same folder:

```
"classes": [
  {"name": "stressinv::Serie",      "header": "stressinv/Serie.h"},
  {"name": "stressinv::Data",      "header": "stressinv/Data.h"},
  {"name": "stressinv::CostMetric", "header": "stressinv/cost.h"}
]
```

factor them out:

```
"namespace": "stressinv",
"header_dir": "stressinv",
"classes": [
  {"header": "Serie.h"},
  {"header": "Data.h"},
  {"name": "CostMetric", "header": "cost.h"}
]
```

(`Serie` and `Data` also drop their `name`, which defaults from the header stem — and the derived name is namespace-qualified too.)

The rules apply identically to `classes`, `functions` and `extensions`:

- `namespace` qualifies an entry name **only when it carries no `::` of its own**. `"Serie"` becomes `stressinv::Serie`; a name containing `::` passes verbatim. So fully qualified spellings, nested classes (`stressinv::Model::Inner`) and sub-namespaces keep working unchanged, and mixing shortened and full spellings in one manifest is fine.
- A **leading `::`** pins an entry to the global namespace: `"::c_entry"` becomes `c_entry`. This is the escape hatch for the odd global function (an `extern "C"` one, say) in an otherwise namespaced manifest.
- `header_dir` is prepended to every entry header, inserting a `/` if missing: `"Serie.h"` becomes `stressinv/Serie.h`. A header living elsewhere can step out relative to it (`"../other/x.h"`), or you can leave `header_dir` unset and spell every path in full.

4.7.1 Grouped entries

One pair of top-level defaults cannot cover a library whose headers live in several folders, or that uses sub-namespaces. For that, an element of `classes` or `functions` may be a **group**: an object carrying `entries` plus its own local defaults, instead of being an entry itself.

```
"namespace": "arch",
"header_dir": "Arch",
"classes": [
  {"name": "Vector3", "header": "math/math.h"},

  {"header_dir": "solvers", "entries": [
    {"name": "GmresSolver", "header": "Gmres.h"},
    {"header": "ParallelSolver.h"}
  ]},

  {"header_dir": "algos", "entries": [
    {"header": "DataSuperposition.h"},
    {"namespace": "sinv", "header_dir": "stress-inversion", "entries": [
      {"header": "JointData.h"},
      {"header": "types.h", "entries": [
        {"name": "MCMCConfig"},
        {"name": "MCMCResult"}
      ]}
    ]}
  ]}
]
```

- A group's `header_dir` **appends below** the inherited one: `Arch + solvers` gives `Arch/solvers/Gmres.h`.
- A group's `namespace` **appends to** the inherited one: `arch + sinv` gives `arch::sinv::JointData`. A leading `::` makes it absolute instead of appending.
- A group may set a `header`: the default header for entries that spell none. That is the natural shape for a run of classes declared by one header — `types.h` above yields `arch::sinv::MCMCConfig` and `arch::sinv::MCMCResult`, both from `Arch/algos/stress-inversion/types.h`. Such entries need an explicit `name`, since the stem fallback would give every one the same name.
- Groups **nest, mix freely with plain entries** in the same array, and work identically under `functions` — where a shared-header group reads especially well (a dozen shape generators all declared by `shapes.h`).
- A group cannot carry a `name`, and `//-comment` keys are ignored on groups as everywhere else.

Everything is resolved at load time, so backends and generated output are byte-identical to the fully spelled form.

4.8 Multiple include directories

When your headers do not all live under one root, give `user_include` an array:

```
"user_include": [ "../geom", "../shared/include", "/opt/thirdparty/include" ]
```

Each entry follows the same resolution rules as the single-string form. Every directory is added to the generated bindings' include path, so a class header is resolved against *all* of them —

the first directory containing the file wins, exactly like a compiler's `-I` search order. The array must not be empty; the single-string form is just the one-directory shorthand.

4.9 A field map

The rest of Part II is organised by what you are trying to do, rather than alphabetically. Appendix A is the alphabetical version.

You want to...	See
say which classes and functions to bind	Chapter 5
choose the languages, and name the modules	Chapter 6
get the generated projects to actually compile and link	Chapter 7
make a foreign container or an Eigen type cross the boundary	Chapter 8
reuse a fixed binding surface across projects	Chapter 9
add docstrings, ranges, UI hints	Chapter 10

Chapter 5

Declaring what to bind

5.1 Classes

`classes` is an array of per-class entries. Each binds one C++ type — a `class`, a `struct`, or an `enum`.

```
"classes": [
  { "name": "Model", "header": "Model.h", "doc": "the model class" },
  { "header": "Point.h" },
  { "name": "space::Vector3", "header": "Vector3.h",
    "annotations": "Vector3.ann.json" }
]
```

Field	Req.	Meaning
<code>header</code>	•	Filename emitted into <code>#include "..."</code> , resolved against <code>user_include</code> . May be a glob (Section 5.2).
<code>name</code>		C++ type name, reachable after including <code>header</code> . May be namespace-qualified. Defaults to the header's stem.
<code>expose</code>		The <i>binding</i> name — what scripts see. Defaults to the unqualified <code>name</code> . See Section 5.3.
<code>annotations</code>		Path (relative to the manifest) to an out-of-line annotation side-car. See Section 10.4.
<code>doc</code>		A description string, used by the document backends.
<code>extensions</code>		Free functions exposed as instance methods. See Section 5.5.
<code>final</code>		Treat the class as non-overridable from the host language: no trampoline is generated even when it has public virtual methods.

Inheritance, multiple constructors, enums, nested user types and `std::vector` members are discovered automatically. You do not list base classes — list the most-derived type you want bound and the bases are flattened into it.

5.1.1 `final`: turning off trampolines

By default, a class with public virtual methods gets a *trampoline* subclass in the backends that support host-language overriding (pybind11, N-API), so a Python or JavaScript subclass can override a C++ virtual. Marking the entry `"final": true` suppresses that: the virtuals still bind as ordinary callable methods, they just cannot be overridden from outside.

It has a second, less obvious use. A class with virtuals is wrapped *as its trampoline subclass*, and a Node member-object property (`mesh.vertices`) stores a `T*` — which requires the wrapped type to actually be `T`, not `Js_T`. `final` is what makes such a class eligible.

5.1.2 Members that will not bind

Some members cannot be compiled into a binding at all: a public field whose type is a non-copyable class, a method returning a reference to such a type, a by-value parameter of one. These are **skipped**, not fatal. The class still binds — as an opaque handle plus whatever members do marshal — and every skip is recorded in `coverage.json` with a reason.

Extension methods (Section 5.5) are how you fill the gaps.

5.2 Binding a whole folder

For the common library shape — one class per header, named after its file — an entry whose `header` is a **shell glob** stands for every header it matches: one class entry per file, its `name` derived from the file's stem.

```
"classes": [
  { "header": "shapes/*.h" },
  { "header": "solvers/**/*.h",
    "exclude": ["solvers/detail/*.h", "solvers/*_impl.h"] }
]
```

The pattern syntax is the same one `user_sources` uses: `*`, `?` and `[...]` within a path component, and a component of exactly `**` matching zero or more directories.

Field	Req.	Meaning
<code>header</code>	•	The glob, resolved against every <code>user_include</code> root with the composed <code>header_dir</code> in front.
<code>exclude</code>		A pattern, or a list of them, whose matches are dropped: headers declaring no bound class, a <code>detail/</code> subtree, an <code>_impl</code> convention.
<code>final</code>		Applies to every matched class — it is a uniform property.

`name`, `expose`, `annotations` and `extensions` are **rejected** on a glob entry: they are per-class and cannot mean anything for a folder.

The rules:

- Only `.h`, `.hh`, `.hpp` and `.hxx` are taken, so `**/*` cannot drag in a `.cpp` or a `README`.
- The derived name is the file's **stem**, qualified by the `namespace` in scope exactly like a spelled-out entry with no `name`. A glob works inside a group and takes that group's `header_dir` and `namespace`.
- A header whose stem is not a C++ identifier (`my-utils.h`) is skipped with a warning — there is nothing to derive a name from.
- A pattern matching nothing is a warning, not an error.

5.2.1 The stem is the class name

The loader never opens a bound header — it lists files. So a glob's only claim about a header is the one its *name* makes, and that claim has to hold exactly:

Point.h must declare a class named *Point*, spelled identically, in the namespace in scope for the entry.

Nothing checks this at load time. A stem naming no such class surfaces as a *compile error in the generated driver* — no member named 'vec3' in namespace 'geom' — which is why the cases below are worth knowing before pointing a glob at a folder.

Your header tree	What the glob does	What to write
Point.h declares Point	binds Point	{"header": "*.h"} — nothing else
Point.h declares Point and Aabb	binds Point only — a file yields one entry	list both explicitly (see the take-over rule)
types.h declares Tolerance, Mode — no types	would bind a non-existent types	list them; the take-over rule keeps the glob off the file
vec3.h declares Vec3	would bind a non-existent vec3	list it, or "exclude": ["vec3.h"]
Point.h declares geom::detail::Point	binds geom::Point — the glob applies the namespace <i>in scope</i>	list it, or put the glob in a group with "namespace": "detail"

The take-over rule

An entry that names a class of a header **speaks for that whole header**: the glob then leaves the file alone. That is what makes the *types.h* and *vec3.h* rows work — you name the classes, and no bogus stem entry appears beside them.

Globbered entries land *after* everything spelled out and yield to it, both by name and by header. So this composes as you would want:

```
"header_dir": "geom",
"namespace": "geom",
"classes": [
  { "header": "*.h", "exclude": ["fwd.h"] },
  { "header": "Mesh.h", "annotations": "Mesh.ann.json" },
  { "header": "types.h", "name": "Tolerance" }
]
```

Every *geom/*.h* but *geom/fwd.h* binds, with *geom/Mesh.h* taking its side-car and *geom/types.h* binding *geom::Tolerance* rather than a non-existent *geom::types*.

The cost is the second row of the table: if a header declares *both* a class named after the file and another one you list, taking the file over drops the first. *rosetta_gen* says so rather than leaving it to be noticed:

```
rosetta_gen: note: "Point.h" is listed explicitly (class "Aabb"), so the pattern
that also matched it contributes nothing there -- class "Point" is NOT bound.
List it too if the header declares both; name the header in the pattern's
"exclude" if it does not (silences this)
```

The note fires on the *types.h* and *vec3.h* rows too, where losing the stem class is exactly what you wanted — add those headers to *exclude* to say so once and quiet it. A name already taken is *not* reported: a side-car entry beside a glob is the intended composition, not a loss.

NOTE

When stems and class names disagree often, `rosetta_gen -init <src_dir>` reads the headers (a heuristic token scan) instead of trusting their names, and writes what it found into a manifest — once, into a file you then edit.

5.3 Renaming: `expose`

Every class binds under **one module-level name**: its `expose` override, or its unqualified C++ identifier. Two entries resolving to the same name is an error — the module attribute and the generated trampoline (`Py_<name>` / `Js_<name>`) would collide.

That happens as soon as two bound namespaces declare the same identifier — a slip-inversion `arch::Data` next to a stress-inversion `arch::sinv::Data`. Rename one side:

```
"classes": [
  {"/": "stress inversion", "namespace": "sinv",
    "header_dir": "stress-inversion", "entries": [
      {"header": "Data.h"},
      {"header": "GpsData.h"}
    ]},
  {"/": "slip inversion", "header_dir": "slip-inversion", "entries": [
    {"header": "Data.h", "expose": "SlipData"},
    {"header": "GpsData.h", "expose": "SlipGpsData"}
  ]}
]
```

Python then sees `arch.Data` (the `sinv` one) and `arch.SlipData`; C++ is untouched. Cross-references follow the rename automatically — a TypeScript signature touching `arch::Data` says `SlipData` — and the emitted C++ spells every bound class by its *qualified* name, so the shared identifier never becomes ambiguous in the generated translation unit.

`expose` is honoured by every backend: the runtime ones, the C-ABI ones (where it also names the generated `.cs` / `.java` file), the UI inspectors, `rest` and `openapi` (route paths and schema names), and the text outputs. It works on an **enum** entry, on a **free function**, and on an **extension method**. It must be a plain identifier.

5.4 Free functions

`functions` binds free (non-member) functions without editing your headers. A namespace has no member list to enumerate, so unlike class members these must be named one by one.

```
"functions": [
  { "name": "transform", "header": "common.h",
    "doc": "Scale and swizzle a point into (x*2, z*3, y*4)" }
]
```

Field	Req.	Meaning
<code>name</code>	•	Function name, optionally namespace-qualified.
<code>header</code>	•	Header declaring it.
<code>doc</code>		Description — free functions carry no in-source annotations.
<code>expose</code>		The binding name. Same rule as a class's.
<code>signature</code>		The C++ function type of the one overload to bind.

Free functions share the module namespace with classes, so `arch::solve` and `arch::sinv::solve` — or a function and a class of the same name — collide exactly like two classes do, and are fixed the same way:

```
"functions": [
  { "name": "arch::solve",      "header": "slip.h" },
  { "name": "arch::sinv::solve", "header": "stress.h", "expose": "solve_stress" }
]
```

5.4.1 Binding one overload: signature

An overloaded free function has no reflection to name: `^^GEO::mesh_union` is ill-formed the moment the namespace declares the name twice, so the entry could not be written at all. Give it the *signature* of the one you want:

```
"functions": [
  { "name": "GEO::mesh_union",
    "header": "geogram/mesh/mesh_surface_intersection.h",
    "signature": "void(GEO::Mesh&, const GEO::Mesh&, const GEO::Mesh&, bool)" }
]
```

The signature is a C++ function type — return type first, then the parameter list — spliced into the generated driver **verbatim**. So it must resolve there: spell the types the way the header does, qualified. `rosetta_gen` checks only its shape (a return type, balanced parentheses); the types are the compiler’s business, and an error names your own spelling back to you.

It is **not a filter** over the overload set. It names one member, and only that one binds. List the entry twice, with two signatures and two `expose` names, to bind two of them:

```
{ "name": "GEO::mesh_union", "header": "...", "expose": "mesh_union_flags",
  "signature": "void(GEO::Mesh&, const GEO::Mesh&, const GEO::Mesh&,
                    GEO::MeshBooleanOperationFlags)" }
```

Backends form the function pointer through a disambiguating `static_cast`, so every backend that emits a pointer binds the selected overload. The one that spells the function’s *reflection* instead — `rest`, which emits `bind_rest_function<^^name>` — skips the entry with a note on `stderr`: there is no reflection for one member of an overload set, and a skipped function is honest where an ambiguous one would not compile.

NOTE

Overload selection is a `functions` feature. An extension method cannot take a `signature` — an extension reaches the backends as a *method*, whose emitters spell its address in many more places than a free function’s single address-of.

5.5 Extension methods

Some libraries keep their real API where no binding generator can reach it: `GEO::Mesh`’s geometry lives behind public member objects with raw `double*` accessors, its I/O helpers are overloaded, its UV coordinates sit in an attribute template.

Rather than hand-write a wrapper *class*, list plain free functions — whose **first parameter is `Cls&`** (or `const Cls&`) — as **extensions** of the bound class. They appear to every backend as ordinary instance methods:

```
"classes": [{
  "name": "GEO::Mesh", "header": "geogram/mesh/mesh.h",
  "extensions": [
    { "name": "georo::set_surface", "header": "mesh_ext.h",
      "doc": "Set vertices + triangles from flat arrays." },
    { "name": "georo::vertices", "header": "mesh_ext.h",
      "doc": "Vertex coordinates as a flat array." }
  ]
}]
```

```
m = geogram.Mesh()
m.set_surface(coords, triangles) # calls georo::set_surface(m, ...)
print(len(m.vertices()) // 3)
```

The receiver is dropped from the exposed signature; the remaining parameters and the return type marshal exactly like a free function's. The method name is the function's unqualified identifier, unless `expose` says otherwise.

Supported by `python`, `nanobind`, `node` and `wasm`, plus all the text backends. Backends that can only emit member pointers — `qt`, `qml`, `csharp`, `java` — skip them, and say so in `coverage.json`.

This is also the escape hatch for an overload a name-keyed target dropped: wrap it in an extension under a distinct name and it comes back.

5.6 Out-parameters

```
bool get_doubles(const std::string &name, vector<double> &out, index_t &dim) const;
```

Three results, C++-style. No host language has that shape, and every backend skipped such a member — the argument its runtime converts is a temporary, which cannot bind to a non-const reference. Name the outputs and they leave the exposed signature and join the return instead:

```
"out_params": {
  "AttributesManager::get_doubles": [1, 2]
}
```

Keys are `"Class::method"` or `"ns::function"`; values are **0-based parameter indices**. A key matching nothing is reported on `stderr` rather than ignored.

```
ok, uv, dim = mesh.facet_corners.attributes().get_doubles("tex_coord")
```

Backend	Shape
python / nanobind	a tuple — (ok, uv, dim)
lua	Lua's own multiple returns
node	an array — <code>const [ok, uv, dim] = ...</code>
wasm	an array too, built as an <code>emscripten::val</code>
typescript	a tuple type: <code>[boolean, number[], number]</code>

The return value comes first when the function has one; with a void return the outputs are all there is.

Never inferred, and that is the design

These two are the same C++:

```
void assign_points(vector<double> &pts, index_t dim, bool steal); // an INPUT
bool get_doubles(const std::string &, vector<double> &out, index_t &dim); // an OUTPUT
```

Guessing wrong silently drops an argument in one direction or a result in the other, so rosetta does not guess. An unmarked mutable reference keeps binding exactly as before — input-only for a container, skipped for a scalar.

Marking a `const` reference does nothing: the callee cannot write through it. A mutable reference to a **bound class** is not an out-parameter either — it already crosses as a handle the callee writes through, which is the more faithful reading and needs no adapter.

5.7 Module init

A library is more than its API. geogram wants `GEO::initialize()`, a run of `CmdLine::import_arg_group()` calls and a C-function-pointer registration to have happened before anything works — none of which is expressible as “bind this name”. `module_init` names statements the generated module runs **when it loads**:

```
"module_init": {
  "headers": ["geogram/basic/command_line.h",
             "geogram/basic/command_line_args.h"],
  "statements": [
    "GEO::initialize(GEO::GEOGRAM_INSTALL_NONE)",
    "GEO::CmdLine::import_arg_group(\"standard\")",
    "GEO::CmdLine::import_arg_group(\"algo\")"
  ]
}
```

A bare array is shorthand for `statements` alone. The statements are spliced **verbatim** into the module entry point — `PYBIND11_MODULE`, `NB_MODULE`, `Init`, `EMSCRIPTEN_BINDINGS`, `luaopen_*`, `JLCXX_MODULE` — at the very top, before any binding is registered. So an I/O handler an init call installs is in place before a bound loader can be reached. A trailing `;` is optional and never doubled.

`headers` are emitted as `#include` lines *ahead* of the bound classes’ own, since an init statement usually names a namespace the bound headers never mention.

| WATCH OUT

Because the statements are spliced verbatim, they must compile at that point: spell types qualified, and remember the module has no arguments to offer. This is *fixed* start-up work, not a configurable entry point. Anything a caller should be able to vary — a verbosity flag, a thread count — still belongs in a bound function.

Emitted by `python`, `nanobind`, `node`, `wasm`, `lua` and `julia` — the backends with a module entry point. The C#, Java, REST, UI and document backends have no single load-time hook and ignore the field.

Chapter 6

Targets

`targets` lists the language backends to emit. One generator produces a **single combined module per target**, exposing every class — not one module per class.

6.1 The two entry forms

A **bare string** uses `module_name` (falling back to `generator_name`) as the module name:

```
"targets": ["python", "node", "markdown"]
```

An **object** sets a per-target module name, and carries the per-target options:

```
"targets": [
  { "lang": "python", "name": "pygeom" },
  { "lang": "node",   "name": "jsgeom" },
  { "lang": "markdown" }
]
```

`name` is optional in the object form too. Mixing the forms in one array is fine.

6.2 The available `lang` values

Every compiled target is **expanded**: reflection runs once on the generation host and the emitted code builds with an off-the-shelf compiler. `rest` and `json` are the two exceptions — both emit a project whose sources include a runtime visitor that still splices reflections, so those two need the C++26 toolchain to build.

lang	Output	Builds with
python	pybind11 extension module	off-the-shelf compiler
nanobind	nanobind extension module	off-the-shelf compiler
node	N-API native addon	off-the-shelf compiler
wasm	Emscripten / embind module	off-the-shelf emsdk
qt	Qt Widgets property/method inspector	off-the-shelf compiler + Qt 6
qml	QML / QtQuick inspector	off-the-shelf compiler + Qt 6
imgui	Dear ImGui inspector app (GLFW + OpenGL3, auto-fetched)	off-the-shelf compiler
csharp	C-ABI shared lib + P/Invoke wrappers	off-the-shelf compiler + .NET

lang	Output	Builds with
java	C-ABI + handle-backed FFM wrappers	off-the-shelf compiler + JDK
julia	CxxWrap.jl / jlccxx shared module	off-the-shelf compiler + CxxWrap.jl
lua	sol2 shared module, require -able	off-the-shelf compiler
rest	cpp-httpplib JSON server + browser client	C++26 toolchain
openapi	OpenAPI 3.1 specification	text output
json	nlohmann (de)serialisation + demo.cpp	C++26 toolchain
typescript	ambient .d.ts declarations	text output
markdown	API reference document	text output
html	styled API reference page	text output
paraview	ParaView Server Manager plugin XML	text output
dynamic	the IR as runtime data (Chapter 13)	off-the-shelf compiler

Notes worth having:

- **Qt / QML** need Qt 6, but no moc on the generated code.
- **Lua** needs Lua 5.1–5.4 or LuaJIT — sol2 does not support Lua 5.5 yet. sol2 itself is fetched at configure time.
- The text-only outputs (markdown, html, typescript, openapi, paraview) compile nothing, so the C++26-vs-off-the-shelf distinction does not apply to them. json is *not* one of them: besides the serialisation header it emits a buildable **demo.cpp** at CMAKE_CXX_STANDARD 26.
- Each name also accepts a deprecated **-expanded** spelling (lua-expanded, qt-expanded, ...) resolving to the same backend, so manifests written before the suffix was dropped keep working.

NOTE

Until 2026-08, seven languages shipped a second, *thin* backend whose generated code re-ran the reflection walk at the target's compile time. Those are gone: the short name now means what **-expanded** used to. No target spells **-expanded** any more, and every old spelling still resolves.

6.3 Per-target linker flags

link_options applies extra linker flags to *this target only*, emitted as target_link_options(... PRIVATE ...) in the generated project.

```
"targets": [
  { "lang": "python" },
  { "lang": "wasm", "link_options": ["-lnodefs.js"] }
]
```

It is per-target — unlike compile_definitions, which is global — because link flags are toolchain-specific. geogram's GEO::initialize() mounts the host filesystem with NODEFS

under Node, which needs emscripten's nodefs library on the **wasm** link line and would break a native link.

The flags are emitted *after* the template's own `target_link_options`, so on the wasm targets — where emcc honours the *last* occurrence of a repeated `-s` setting — they can also override a template default, e.g. `"-sALLOW_MEMORY_GROWTH=0"` for a fixed-size heap.

6.4 Where the built artifact lands

A generated project drops its built module next to its own sources. Handy for a smoke test; useless when the `.so` belongs inside your Python package or the `.js` next to a web app's assets.

```
"targets": [
  { "lang": "nanobind", "name": "implicit3d", "out_dir": "./dist" },
  { "lang": "wasm",      "name": "implicit3d", "out_dir": "../www/assets" }
],
"out_dir": "./dist"
```

The top-level `out_dir` is the default for every target that names none; a target's own entry wins. Paths resolve against the manifest's directory and are created if missing.

WATCH OUT

This is **not** where the generated project goes — that is `rosetta_gen`'s own output tree, controlled by `-bindings-dir`. It is where the *loadable artifact* lands: `libfoo.so` / `foo.pyd`, `foo.node`, `foo.so` for Lua, and for wasm *both* halves, the `.js` loader and its `.wasm`.

The copy is `copy_if_different`, so an unchanged build touches nothing downstream, and the existing next-to-the-sources copy still happens.

Supported by every backend that builds a loadable module: `python`, `nanobind`, `node`, `wasm`, `lua`, `julia`. The document backends have no artifact to place, and the application-shaped ones (`qt`, `qml`, `ingui`, `rest`, `csharp`, `java`) are not covered.

6.5 Pinning the runtime

By default a generated project takes whatever the `PATH` gives it: the emitted CMake probes `python3` / `node`, with a 3.8 floor and N-API 8 written in. Pin them per target instead:

```
"targets": [
  { "lang": "nanobind", "name": "geom",
    "python": "3.11", "requires_python": ">=3.10" },
  { "lang": "node", "name": "geom",
    "napi_version": 9, "node_engine": ">=18" }
]
```


Field	Effect	Default
<code>python</code>	The interpreter the binding is built for. A bare version ("3.11") becomes <code>python3.11</code> , resolved on PATH; anything else is used as written (" <code>/opt/py311/bin/python3</code> ").	probe <code>python3</code>
<code>requires_python</code>	Minimum version. Feeds both <code>find_package(Python ...)</code> and <code>pyproject.toml</code> 's <code>requires-python</code> , so the build floor and the wheel metadata cannot drift apart.	<code>>=3.8</code>
<code>napi_version</code>	<code>NAPI_VERSION=</code> for the addon. This is the real Node floor — N-API 8 means Node 12.22+, N-API 9 means Node 18.17+.	8
<code>node_engine</code>	<code>engines.node</code> in <code>package.json</code> — documentation for npm rather than a compile setting, which is why it is separate from <code>napi_version</code> .	absent

These are per-target, not top-level, so a manifest carrying both a `python` and a `nanobind` target can point them at different interpreters.

Two things worth knowing:

- `-cmake-arg -DPython_EXECUTABLE=...` **does not work**, which is why these fields exist. The generated CMake sets that variable with `CACHE ... FORCE` after probing, overriding anything passed on the command line. Before these fields, a `venv` or a `pyenv` shim was the only way to choose an interpreter.
- **python also drives the wheel.** A wheel is tagged for whichever interpreter runs `make_wheel.py`, so `rosetta_gen -build -wheel` packages with the pinned interpreter rather than the probed one — otherwise a 3.11 build would emit a `cp312` wheel.

6.6 Choosing targets

A few rules of thumb, since nineteen is a lot:

Python

`python` (`pybind11`) is the safe default — the most permissive type gate, full overload sets, mature. `nanobind` is leaner and faster and produces `abi3` wheels on 3.12+, at the cost of a slightly narrower feature set. Emitting both costs nothing but build time.

JavaScript

`node` for a native addon, `wasm` for the browser. Add `typescript` in either case: it costs nothing to generate and describes the Node addon's surface to an editor.

Documentation

`markdown` or `html` are free — no compilation, no toolchain, and they stay in step with the code by construction. There is little reason not to emit one.

A user interface

`qt`, `qml` and `imgui` each generate a working inspector over your annotated types. If you want a UI that is a *query* rather than generated code, see Chapter 13 instead.

During development

Generate everything, but compile one or two: `-only python,node` (Section 11.4.2).

Chapter 7

Making the generated projects build

Chapter 5 said what to bind. This chapter is about the other half: getting the emitted projects to actually compile and link, which for a real library means telling rosetta where the *definitions* are, what preprocessor state the headers expect, and which optimisation to build at.

7.1 The two ways to supply definitions

Your bound headers usually only *declare* the API. The bodies live either in a pre-built library or in source files. Rosetta has one field for each, and they are independent — use either, or both.

<code>user_lib</code>	Link the bindings against a pre-built <code>.so</code> / <code>.dylib</code> / <code>.a</code> .
<code>user_sources</code>	Compile <code>.cpp</code> (or <code>.c</code>) files <i>straight into</i> each binding target.

7.1.1 Linking a pre-built library

```
"user_lib": {
  "name": "space",
  "dir": "../space/bin",
  "link": "shared"
}
```

Field	Req.	Meaning
<code>name</code>	•	Library name — the <code>space</code> in <code>libspace.dylib</code> .
<code>dir</code>	•	Directory holding it, relative to the manifest; used for <code>-L</code> and <code>rpath</code> .
<code>link</code>		"shared" (default), "static", or "dynamic" (an alias of shared). The <i>preferred</i> form, with fallback to whichever is actually built.

A bound library rarely stands alone. Give `user_lib` an **array**, one object per library:

```
"user_include": ["../mylib/include", "/opt/foo/include"],
"user_lib": [
  { "name": "mylib", "dir": "../mylib/bin", "link": "shared" },
  { "name": "foo", "dir": "/opt/foo/lib" },
  { "name": "bar", "dir": "/opt/foo/lib", "link": "static" }
]
```

- **Order is the link order** — list a library before the ones it depends on, which is what static archives require.
- Each entry keeps its own `link`, so a shared library and a static dependency mix freely.
- Every distinct `dir` lands in the binding's `BUILD_RPATH` and `INSTALL_RPATH`, so the module loads without `LD_LIBRARY_PATH` / `DYLD_LIBRARY_PATH`.
- The **generator driver** links the same list — it instantiates the bound types during the reflection walk, so it needs the dependencies' definitions too, not just yours.
- Headers of the dependencies go in `user_include`; their configuration macros in `compile_definitions`.

WATCH OUT

On `wasm`, every entry is taken as the static archive `lib<name>.a` built with the **same** `emsdk`. A native `.dylib` / `.so` cannot enter a `wasm` module, so a dependency you cannot rebuild for `emscripten` rules that target out.

Flags that are not a `-L/-l` pair — a framework, a `-start-group`, a package's own link line — go through a target's `link_options` (Chapter 6); the two compose.

7.1.2 Compiling sources in

```
"user_sources": [
  "../src/widget.cpp",
  "../src/shape.cpp"
]
```

Always a list of paths, each relative to the manifest (or absolute); a single string is accepted as a one-element list. Every compiled backend adds them via `target_sources(...)`, so they build with the same include path and flags as the generated binding.

Entries may be **shell globs**, expanded at generation time:

```
"user_sources": ["extern/arch/src/**/*.*c"]
```

`*`, `?` and `[...]` work within a path component, and a component that is exactly `**` matches **zero or more directories** — so one pattern covers a whole source tree. Like the other wildcards, `**` does not enter hidden directories. Matches are sorted for reproducible output and the final list is de-duplicated, so mixing a literal path with a glob that also covers it is safe. A pattern matching nothing warns and is skipped.

WATCH OUT

Globs expand when `rosetta_gen` runs. **Re-run it after adding or removing matching files** — a new `.cpp` does not appear in the build by itself.

Entries may also be **C sources**. When any `.c` file is listed, the generated `CMakeLists.txt` calls `enable_language(C)` automatically, which is what makes a library's vendored dependencies work:

```
"user_sources": [
  "../geogram/src/lib/geogram/mesh/*.cpp",
  "../geogram/src/lib/geogram/third_party/rply/*.c",
]
```

```

    "/geogram/src/lib/geogram/third_party/zlib/*.c"
]

```

The text-only backends compile nothing and ignore the field.

7.2 Preprocessor definitions

Most often a third-party library's configuration macros. Each entry is "NAME" or "NAME=VALUE":

```

"compile_definitions": [
    "GEOGRAM_USE_BUILTIN_DEPS",
    "GEOGRAM_WITH_HLBFGS"
]

```

They are emitted as `target_compile_definitions(... PRIVATE ...)` in **two** places, so the whole pipeline sees a consistent configuration:

- the **generator driver** — it includes the bound headers for the reflection walk, which must see the same preprocessor state the bindings will be built with;
- **every compiled binding target**, where they reach the bound headers and the `user_` sources alike.

That first bullet is the one that matters. A macro that changes a class's member set — and plenty do — would otherwise produce a binding for a different class than the one you compile.

A single string is accepted as a one-element list.

7.3 Build type and optimisation

```

"build_type": "Release",
"optimization": "-O2"

```

Both optional and independent.

`build_type` is one of `Debug`, `Release`, `RelWithDebInfo`, `MinSizeRel` (case-insensitive). It is emitted as a *default* inside `if(NOT CMAKE_BUILD_TYPE AND NOT CMAKE_CONFIGURATION_TYPES)`, so `-DCMAKE_BUILD_TYPE=...` at configure time still wins and multi-config generators (Visual Studio, Xcode) are left to pick a configuration at build time.

NOTE

Omitted means Release. It used to mean “emit nothing”, which left CMake with no build type and hence no optimisation flags at all — so the documented two-line build produced an unoptimised module, around eight times slower on a trivial pybind11 call. Pass `-DCMAKE_BUILD_TYPE=Debug`, or set `"build_type": "Debug"`, when you want an unoptimised build.

`optimization` is an explicit level: `-O0` through `-O3`, `-Os`, `-Oz`, `-Og` or `-Ofast` (the leading `-` may be omitted). It is emitted as `add_compile_options(...)` and `add_link_options(...)`, which land *after* the build type's own per-config flags — so this `-O` overrides the level the build type implies. `"build_type": "Release", "optimization": "-O2"` builds `-DNDEBUG` but at `-O2` instead of Release's `-O3`. The link option matters for the wasm targets, where emscripten optimises at link time too.

Both apply to the bindings and to any `user_sources` compiled into them, in every compiled backend.

7.4 Headers the library's own build produces

A header configured by the bound library's build system is simply absent when rosetta compiles that library's sources without running its build. geogram's `<geogram/version.h>` is the case in point — configured by geogram's CMake from `version.h.in`.

```
"generated_headers": [
{
  "path": "geogram/version.h",
  "template": "../extern/geogram/src/lib/geogram/basic/version.h.in",
  "substitutions": {
    "VORPALINE_VERSION_MAJOR": "1",
    "VORPALINE_VERSION_MINOR": "10",
    "VORPALINE_VERSION_PATCH": "1",
    "VORPALINE_VERSION": "1.10.1-rosetta",
    "VORPALINE_BUILD_NUMBER": "",
    "VORPALINE_BUILD_DATE": "",
    "VORPALINE_SVN_REVISION": ""
  }
}
]
```

Field	Req.	Meaning
<code>path</code>	•	The <i>relative include path</i> — exactly what the sources <code>#include</code> .
<code>template</code>	one of	A file with <code>@KEY@</code> placeholders (CMake's <code>configure_file</code> form), resolved against the manifest's directory.
<code>content</code>	one of	Literal lines, as an array of strings, for a header with no template.
<code>substitutions</code>		<code>@KEY@</code> to value, applied to <code>template</code> .

The file is written to `<bindings>/include/<path>` and that directory is placed **first** on every generated target's include path — ahead of the library's own sources, so a stale checked-in copy of the same header cannot win the lookup.

Substitution happens when the manifest is read, so a missing key is reported against your manifest (no substitution for `@VORPALINE_VERSION@`) rather than compiling a literal `@KEY@` into the binding. Only the `@KEY@` form is recognised, since `${KEY}` would collide with the shell and CMake expansions such templates carry.

NOTE

A `-D` define is a different tool for a nearby job. geogram's `GEOGRAM_VERSION` is passed by its CMake on the command line, not written into the header, so it belongs in `compile_definitions`.

7.5 Python wheels

The `python` and `nanobind` backends emit, alongside `CMakeLists.txt`, everything needed to build a redistributable wheel:

<code>pyproject.toml</code>	scikit-build-core build config. It drives the generated <code>CMakeLists.txt</code> itself, so the wheel and a plain <code>cmake -build</code> compile the same target the same way — there is no second build description to keep in sync.
<code>make_wheel.py</code>	One-command builder: builds the wheel, then bundles external shared libraries and fixes the platform tag. Python rather than a shell script, so Linux, macOS and Windows are covered by one file.

```
cd bindings/nanobind
python make_wheel.py           # -> dist/*.whl, dist/repaired/*.whl
python3.11 make_wheel.py       # a specific interpreter
python make_wheel.py --outdir /tmp/whl # somewhere other than ./dist
python make_wheel.py --no-repair # skip the bundling / retagging step
```

The wheel is built for whichever interpreter runs the script — there is no `-python` flag, just invoke the one you want (or pin it with the target’s `python` field, Section 6.5).

`version` is the distribution version stamped into the packaging artifacts, PEP 440 style (“1.2.0”, “0.3.0rc1”). Default 0.1.0.

Two fields make packaging automatic for a project that always ships wheels. They go **on the target**, beside the interpreter it is built for:

```
"targets": [
  { "lang": "python", "name": "pygeom", "python": "3.11" },
  { "lang": "nanobind", "name": "nbgeom", "wheel": true, "wheel_dir": "./dist" }
]
```

<code>wheel</code>	Build a wheel for <i>this</i> target on every <code>-build</code> , without passing <code>-wheel</code> .
<code>wheel_dir</code>	Where this target’s wheels go instead of its own <code>dist/</code> . Implies <code>wheel</code> , and resolves against the manifest’s directory.

Only the two backends that emit a `make_wheel.py` accept them: `wheel` on a `markdown` or `node` target is an error, not a no-op.

Per-target rather than top-level, and for the same reason `out_dir` is (Section 6.4): the two Python backends genuinely differ here. `nanobind` produces one `abi3` wheel covering every CPython 3.12+ where `python` produces one per version, and `python` (Section 6.5) already pins per target which interpreter each is tagged for. Sharing one `wheel_dir` between them is safe — each `make_wheel.py` repairs only the wheels it just produced.

The **flags still win, in the direction of doing more**. `-wheel` packages every `python` / `nanobind` target whatever the manifest says, so a target spelling `"wheel": false` cannot turn off a run that asked for it; `-wheel-dir DIR` overrides every target’s own directory, which is what makes “collect them all in one place” a single flag. With neither set anywhere, `-build` only compiles the extension module.

WATCH OUT

`wheel` and `wheel_dir` were **top-level** keys. They are per-target now, and a top-level one is a load error naming the fix rather than a key that silently does nothing — a manifest that used to ship wheels must not quietly stop.

NOTE

On 3.12+, nanobind wheels are tagged `abi3`, covering every later CPython with a single artifact.

7.6 A realistic manifest

Putting the last four chapters together — this is roughly the shape a real library’s manifest takes:

```
{
  "//": "Bindings for the geom library.",

  "variables": [
    { "name": "SRC", "value": "./extern/geom" }
  ],

  "user_include": ["$SRC/include", "./extern/eigen-3.4.0"],
  "rosetta_include": "./extern/rosetta/include",
  "generator_name": "geom_gen",
  "module_name": "geom",

  "user_sources": ["$SRC/src/**/*.cpp"],
  "compile_definitions": ["GEOM_NO_ASSERT"],
  "build_type": "Release",
  "version": "1.2.0",

  "interop": ["eigen"],
  "sequences": [{ "type": "Eigen::VectorXd" }],

  "namespace": "geom",
  "header_dir": "geom",

  "targets": [
    { "lang": "python", "name": "pygeom", "requires_python": ">=3.10" },
    { "lang": "nanobind", "name": "nbgeom", "out_dir": "./dist",
      "wheel": true, "wheel_dir": "./dist/wheels" },
    { "lang": "node", "name": "jsgeom" },
    "typescript",
    "markdown"
  ],

  "classes": [
    { "header": "*.h", "exclude": ["fwd.h", "detail/*.h"] },
    { "header": "Mesh.h", "annotations": "Mesh.ann.json" }
  ],

  "functions": [
    { "name": "remesh", "header": "algorithms.h", "doc": "Isotropic remeshing." }
  ]
}
```

Chapter 8

Types across the boundary

Reflection tells rosetta that a method takes a `GEO::vector<double>`. It does not tell it what a Python caller should hand over. That mapping — the *type gate* — is per backend, and this chapter is about widening it for the types your library actually uses.

The default gate handles the obvious cases without a word in the manifest: arithmetic types, `bool`, `std::string`, enums, `std::vector<T>` of any of those, and any *bound class*. Everything below is for the rest.

NOTE

The IR's rule for anything it cannot claim is uniform: mark the type `kind: "unknown"` and let each backend skip the member. That is why a type rosetta does not understand costs you one skipped method with a stated reason, not a compile error.

8.1 Foreign sequence containers

Many libraries carry their bulk data in their *own* vector type — geogram's `GEO::vector<T>`, an aligned or pooled vector — and the marshalling layers only know `std::vector`. List the container template (qualified, **one type parameter**) and it crosses like a `std::vector` of its element:

```
"sequences": ["GEO::vector"]
```

rosetta_gen emits

```
template <typename T>
struct rosetta::is_sequence<GEO::vector<T>> : std::true_type {};
```

into the generated `bindings.h`. The container must be default-constructible with `value_type`, `size()`, `resize(n)` and `begin()/end()`; elements must be arithmetic, `bool`, `std::string` or a bound enum.

The opted-in backends — `python`, `nanobind`, `node`, `wasm`, `lua`, plus `typescript` declarations — marshal it **by copy through a `std::vector<element>` boundary** inside an emitted adapter. Scripts pass and receive plain arrays, lists or tables. Every other backend keeps skipping the type.

Three consequences worth knowing:

- **Mutable Seq& parameters bind, input-only.** The adapter's local container is a real lvalue, so geogram's `assign_points(vector<double>&, dim, steal)` works — `steal` steals from the adapter's copy, which is fine. In-place mutations are discarded, exactly like pybind's `std::vector&` casters.
- **Overload sets need the sequence overload to be first.** The adapter calls the method *by name* with concrete arguments instead of spelling the ambiguous `&T::name` member pointer. A backend that keys methods by name binds the first-declared overload, so a set whose first declaration is the raw-pointer one stays skipped there.
- **Virtual methods naming the container cannot be overridden script-side** — the exact spelling cannot round-trip — but they still bind as callable methods.

Sequence-typed public *fields* bind as copying properties in python, nanobind, wasm and lua; node skips them.

8.1.1 Registering a concrete type

The string form registers a *template*, and the adapter spells the container it builds as `Namespace::Template<element>`. That composition is wrong the moment the specialisation carries more than its element: `Eigen::VectorXd` is really `Eigen::Matrix<double, -1, 1>`, and the composed `Eigen::Matrix<double>` does not compile.

Register the concrete type instead, spelled exactly as the adapter should write it:

```
"sequences": [{ "type": "Eigen::VectorXd" }]
```

which emits a full specialisation — plus the spelling, as `rosetta::is_sequence` and `rosetta::sequence_cpp_name` — rather than the partial one.

`{ "template": "GEO::vector" }` is the long form of a plain string entry. An object with *both* keys, or neither, is an error — the two cannot be told apart by looking at the text, and guessing wrong emits code that does not compile.

This is also the only way to register a container that is not a template at all.

8.2 Foreign matrices

sequences one dimension up. A 2-D type — `Eigen::MatrixXd`, a library's own dense grid — is a sequence in no useful sense, so it needs its own registration:

```
"matrices": [{ "type": "Eigen::MatrixXd" }]
```

Entries take the **same two forms** as sequences: a plain string for a one-type-parameter template (`"mylib::Grid"`, spelled `mylib::Grid<double>` by the adapter), or `{ "type": "..."` } for a concrete type.

The type must be default-constructible with `value_type`, `rows()`, `cols()`, `resize(r, c)` and `operator()(i, j)`; the element must be **arithmetic** — a grid of strings has no natural script shape.

The opted-in backends marshal it **by copy through a `std::vector<std::vector<element>`** boundary — an array of row arrays. typescript says `number[] []`.

- **Row-major, whatever the matrix stores.** `operator()(i, j)` is the only access the registration promises, so the boundary is built row by row regardless of the type's own storage order.

- A **ragged incoming array is squared off** to the first row's length: rows are the outer size, columns the first row's; short rows keep whatever `resize()` left and extra columns are dropped.
- A mutable `Mat&` parameter **binds input-only**, like a sequence one.

8.3 Foreign-library interop

A method taking or returning `Eigen::VectorXd` is, to the reflection walk, a method taking an ordinary class. Every backend used to bind it happily — and the call then **threw at run time**, because no caster for that class was ever registered.

`interop` names foreign libraries whose types the target's binding framework can already marshal on its own:

```
"interop": ["eigen"]
```

Recognition is by **enclosing namespace**, not by a list of type names, so one line covers everything the library owns: `VectorXd`, `MatrixXd`, `Vector3d`, `Map`, `Ref`, `Block`, `Array`, the sparse types — including spellings a hand-maintained list would have missed.

Backend	Behaviour
python	Emits <code>#include <pybind11/eigen.h></code> ; the types bind natively as numpy arrays , both directions, matrices included. No adapter, no copy where the layout allows a view.
nanobind	Same, through <code><nanobind/eigen/dense.h></code> . (Sparse is not emitted — <code><nanobind/eigen/sparse.h></code> costs compile time a rare signature does not justify.)
every other backend	Skips the member. Deliberate, and an improvement: a skipped method is honest, a bound one that always throws is not.

Two practical notes: Eigen must be on the include path of the generated binding (it already is if `user_include` lists it, which it must anyway for the bound headers to compile), and the caster header is only emitted when a bound signature actually names such a type — so declaring `interop` on a library that never exposes one costs nothing.

8.3.1 Reaching the caster-less backends

`node`, `wasm` and `lua` have no caster to lean on, so on their own they skip every Eigen-typed member. Give them the flat array by **also** registering the concrete type as a sequence:

```
"interop": ["eigen"],
"sequences": [{ "type": "Eigen::VectorXd" }]
```

The IR then carries both marks, and each backend picks:

python / nanobind	The caster wins — still numpy, still no copy. A dual-marked type costs them nothing.
node / wasm / lua	Bind through the sequence adapter : in, out and back through a <code>std::vector<double></code> boundary. Scripts see a plain array.
typescript	Declares <code>number[]</code> , matching what those runtimes hand out.

1-D only. `MatrixXd` has no useful reading as a sequence — register it under `matrices` instead, which splits the backends exactly the same way:

```
"interop": ["eigen"],
"sequences": [{ "type": "Eigen::VectorXd" }],
"matrices": [{ "type": "Eigen::MatrixXd" }]
```

covers both shapes, in both directions, on every expanded backend.

8.4 Two standard types that need no declaration

Nothing to write in the manifest for these. They are here because a factory API tends to name both at once, and until recently either one made a signature quietly unbindable on the caster-less backends:

```
std::shared_ptr<Mesh> compile_file(const std::filesystem::path &input);
```

`std::filesystem::path` marshals as a string

It is described as kind: `"string"` with an `is_path` flag.

python / nanobind	Native, through <code><pybind11/stl/filesystem.h></code> / <code><nanobind/stl/filesystem.h></code> — callers pass a <code>str</code> or any <code>os.PathLike</code> , and a returned path comes back as <code>pathlib.Path</code> .
node / wasm / lua	Through the same copy adapter the foreign containers use: the boundary declares <code>std::string</code> , the emitted code builds a path from it and calls <code>.string()</code> on the way back.
typescript	Declares <code>string</code> .

`std::shared_ptr<T>` crosses in the return direction

Which is where a factory needs it. The pointee `T` must itself be a bound class.

python	<code>py::class_<T, std::shared_ptr<T></code> — the holder is declared automatically for every pointee the module hands out.
nanobind	Native, through <code><nanobind/stl/shared_ptr.h></code> .
wasm	The class registration gains <code>.smart_ptr<std::shared_ptr<T>("T_sp")</code> , and only for classes that actually travel that way.
lua	Native — sol2's own <code>unique_usertype_traits</code> handles it.
node	The JS object adopts the pointer: a third ownership mode in the wrapper, next to “owns” and “aliases a member of a pinned parent”. The C++ object lives as long as the JS handle does.
typescript	Declares the <i>pointee</i> (Doc), not <code>shared_ptr</code> .

Two limits:

- **Returns only.** A `shared_ptr<T>` *parameter* stays skipped: converting an incoming script object would mean manufacturing a control block over memory the binding does not own — a double-free waiting to happen. Take `const T&`.
- **node: untrampoline pointees only.** A class with virtual methods is wrapped as its trampoline subclass, and the adopting wrapper stores a `T*`; reading one as the other

is undefined, so such a return stays skipped there. The other backends have no such restriction.

Both work with a non-copyable `T` — which is the point, since a class handed out by a factory usually is one.

8.5 Deciding what you need

Your signature names...	Write
your own <code>Vec<T></code> template	"sequences": ["my::Vec"]
<code>Eigen::VectorXd</code> , Python only	"interop": ["eigen"]
<code>Eigen::VectorXd</code> , everywhere	"interop" + "sequences": [{"type": "Eigen::VectorXd"}]
<code>Eigen::MatrixXd</code> , everywhere	"interop" + "matrices": [{"type": "Eigen::MatrixXd"}]
a 2-D grid of your own	"matrices": ["my::Grid"]
<code>std::filesystem::path</code>	nothing
<code>std::shared_ptr<Bound></code> returned	nothing
a <code>std::function</code> parameter	nothing will help — the member is skipped, with a reason. Add an <code>extensions</code> wrapper if you need it.

When in doubt, generate and read `coverage.json`: it names every type that failed a gate, and which one.

Chapter 9

Presets

Some binding surfaces are **fixed**: they do not depend on your project at all. Binding rosetta's own reflection API is the case that motivated this — the twelve `rosetta::script` handle classes and the eight registry free functions are the same in every project that wants them, and copying twenty entries into each manifest makes the three lines that *are* project-specific hard to find.

9.1 Using one

`preset` names a JSON fragment shipped under

`<rosetta_include>/rosetta/presets/<name>.json`

```
{
  "preset": "scriptable",
  "rosetta_include": "../..//include",
  "user_include": ["../lib", "../lib/bindings/dynamic"],
  "user_sources": ["../lib/bindings/dynamic/auto_dynamic.cpp"],
  "module_init": { "headers": ["auto_dynamic.h"],
                   "statements": ["bank::register_all()"] },
  "targets": ["python", "lua"]
}
```

That is a complete manifest. Six keys, every one of them about *this* project: where its files are, which targets to emit, what to run at load. `classes`, `functions` and `module_name` all come from the preset — and a preset that supplies `module_name` is what lets the targets stay bare strings.

Give an array to apply several: `"preset": ["scriptable", "..."]`, applied in order. An unknown name is an error naming the path that was searched, so a typo surfaces immediately.

9.2 What applying one does

Merge	Deep, with the manifest winning every conflict it can express. A key you did not write is taken from the preset; arrays present in both are concatenated <i>preset first</i> ; objects merge recursively; a scalar you spelled out is never overridden.
classes / functions	Parsed with a pristine context : your namespace and header_dir defaults describe <i>your</i> tree and are not prepended to a preset's headers. They are appended after your own entries.
Include path	A preset binds headers that live under rosetta_include , so it puts that directory on the target's include path too. Every other binding is reflection-free by construction and has no reason to see rosetta's headers, which is why this is not the default.

The array rule is what makes `module_init` compose: the preset contributes its `caster include` and `your` statements survive, rather than one silently replacing the other.

`classes` becomes optional when a preset supplies it — which is why the manifest above has none.

9.3 The shipped preset

Name	What it binds
<code>scriptable</code>	<code><rosetta/script.h></code> — the bindable facade over the runtime object model — plus the per-language Value casters. Lets a host language walk the metadata of any library that emits dynamic tables, and call into it by name.

Chapter 14 is what that preset is *for*, and is worth reading before using it — it is the most unusual thing in rosetta and the hardest to guess at from the manifest alone.

9.4 Writing your own

A preset is only a JSON fragment. Drop one into `<rosetta_include>/rosetta/presets/` and it is available by name to every manifest that points at that include directory.

It is worth doing when a binding surface is genuinely fixed across projects:

- a company-wide utility library every binding also exposes;
- a facade type plus its helpers, where the twenty entries never vary;
- a set of `sequences / matrices / interop` registrations shared by every project using the same numerical stack.

It is *not* worth doing to factor out the parts of a manifest that differ per project — `user_include`, `user_sources`, `targets` — because the merge rule (manifest wins) means those would be overridden anyway, and a preset nobody can use unchanged is just an indirection.

The fields most worth putting in one:

```
{
  "module_name": "mylib",
  "classes": [ ... fixed entries ... ],
  "functions": [ ... fixed entries ... ],
  "module_init": { "headers": ["mylib/init.h"],
```

```
      "statements": ["mylib::initialize()"] },  
    "sequences": ["mylib::Vector"]  
  }
```

NOTE

Remember the pristine-context rule when writing one: a preset's `classes` entries must spell their headers and names in full, because the consuming manifest's `namespace` and `header_dir` will *not* be applied to them.

Part III

Annotations

Chapter 10

Annotations

Annotations are rosetta's **opt-in enrichment layer**: small declarative tags attached to your fields and methods that backends turn into docstrings, validation and user interface. They are never required — reflection alone binds the whole class — and they never change the C++ behaviour of your types.

They buy three things:

1. **Documentation** — `doc` becomes docstrings, reference pages, tooltips.
2. **Validation** — `range` rejects bad assignments at the binding boundary.
3. **User interface** — `label`, `button`, `combobox` and the `widget::*` hints drive the generated Qt, QML and Dear ImGui inspectors.

10.1 Two spellings, one result

Every annotation can be written **inline** (as a P2996/P3394 attribute in the header) or **out of line** (as a JSON side-car, leaving the header plain C++). The two forms produce identical bindings and can be mixed on the same class.

```
#include <rosetta/annotations.h>

class Algo {
public:
    [[= rosetta::doc{"Max solver iterations"}, = rosetta::range{0, 200},
     = rosetta::widget::slider, = rosetta::label{"Max iter"} ]]
    int maxIter{100};

    [[= rosetta::doc{"Run the solver"}, = rosetta::button{"Run"} ]]
    double run();
};
```

...or the same thing with a clean header and a side-car:

```
{
  "maxIter": { "doc": "Max solver iterations", "range": [0, 200],
               "widget": "slider", "label": "Max iter" },
  "run":     { "doc": "Run the solver", "button": "Run" }
}
```

Annotations are read *once*, on the generation host, and merged (inline first, then JSON) into the same annotation list. Backends cannot tell the two sources apart.

10.2 The annotation set

Kind	Applies to	Spellings	Meaning
doc	fields, methods, classes*	<code>= rosetta::doc{"..."}</code> <code>"doc": "..."</code>	Description text. Docstrings, reference pages, OpenAPI descriptions, hover tooltips.
range	numeric fields	<code>= rosetta::range{lo, hi}</code> <code>"range": [lo, hi]</code>	Value bounds. Compiled backends emit a validating setter ; UI inspectors render a clamping slider; doc backends print the bound. Scientific notation is fine (<code>[1e-10, 1e-6]</code>).
readonly	fields	<code>= rosetta::readonly</code> <code>"readonly": true</code>	Getter-only property; writes are rejected per backend; UI editors grey out.
combobox	string fields	<code>= rosetta::combobox>{"a","b"}</code> <code>"combobox": ["a","b"]</code>	Allowed choices (maximum 16). UI renders a drop-down.
label	fields, methods	<code>= rosetta::label{"..."}</code> <code>"label": "..."</code>	Display-name override in UI backends. The C++ identifier stays the lookup/binding key — only the visible text changes.
button	methods	<code>= rosetta::button{"Run"}</code> <code>"button": "Run"</code>	Caption of the call button a UI backend renders for the method.
widget::*	fields	<code>= rosetta::widget::slider</code> <code>"widget": "slider"</code>	Editor hint — see below.

* Class-level and free-function doc come from the **manifest** (a doc key on the entry), since free functions carry no in-source annotations.

10.2.1 Widget hints

Hints pick the editor when more than one would fit the field's type. They are consumed by the three UI inspector backends and ignored everywhere else.

Hint	Field type	Renders as
spin	number	spin box / typed input
slider	number + range	slider with the range as bounds (clamps by construction)
textfield	number	validated free-text input — the right editor for values a slider cannot hit precisely (an eps of <code>1e-7</code>)
checkbox	number (0/1), bool	checkbox
color	string "#rrggbb"	colour picker (ImGui ColorEdit3, Qt swatch + QColorDialog, QML ColorDialog); edits write back as hex
multiline	string	multi-line text area
radio	string + combobox	the choices as a radio-button group instead of a drop-down
file	string path	text entry + browse button opening the platform file dialog

10.3 Who consumes what

Scripting backends

(python, nanobind, node, wasm, lua, julia, plus csharp / java): doc becomes docstrings where the framework has them; range becomes validating setters; readonly becomes getter-only properties. label, button and widget::* are ignored — they are UI concepts.

UI inspectors

(qt, qml, imgui): everything. doc as tooltips, range as slider bounds or commit validation, readonly as disabled editors, combobox as drop-downs, plus label, button and widget::*.

Document backends

(markdown, html, openapi, typescript): doc and the constraints (range, readonly, combobox) are rendered into the reference text or schema.

An annotation a backend has no use for is simply ignored — never an error.

10.4 Out-of-line annotations

The side-car keeps your header plain C++. That matters for three reasons, and the second is the one that usually decides it:

	Inline	Side-car
Header stays plain C++	no	yes
Generated binding builds on an off-the-shelf compiler	no	yes
Works on third-party headers you cannot edit	no	yes
Expressiveness	full set	full set

The middle row is the load-bearing one. Inline annotations pull <experimental/meta> into your header, and the generated binding #includes that header — so the “reflection-free by construction” promise (Section 1.6) is lost for every target.

10.4.1 Wiring it up

One manifest field per class:

```
"classes": [
  { "name": "Algo", "header": "Algo.h", "annotations": "Algo.ann.json" }
]
```

The path is resolved relative to the manifest. The side-car’s keys are member identifiers — fields *or* methods — and each value object may carry any of the annotation kinds:

```
{
  "title": { "doc": "The widget title", "label": "Title" },
  "count": { "doc": "Visible items", "range": [0, 100], "widget": "slider" },
  "id": { "readonly": true },
  "mode": { "combobox": ["fast", "slow"] },
  "run": { "doc": "Run the thing", "button": "Run" }
}
```

The bound header never mentions rosetta:

```
#pragma once
#include <string>
struct Widget {
    std::string title;
    int         count = 0;
    int         id     = 0;
    std::string mode;
};
```

That is the whole surface. Run `rosetta_gen`, build, and every backend sees the annotations.

10.4.2 What it does under the hood

`rosetta_gen` bakes the JSON bytes into `bindings.h`, right after the class header is included so `T` is complete:

```
template <> inline constexpr auto rosetta::detail::ann_storage<Widget> =
    std::to_array<char>({ char(0x7b), /* baked JSON bytes */ '\0' });
template <> inline constexpr std::string_view rosetta::ann_json_source<Widget> =
    std::string_view{ rosetta::detail::ann_storage<Widget>.data(),
                     rosetta::detail::ann_storage<Widget>.size() - 1 };
static_assert(rosetta::detail::ann_keys_error<Widget>().empty(),
              rosetta::detail::ann_keys_error<Widget>());
```

A small consteval parser reads it at `walk()` time and concatenates the entries with the member's inline annotations. Nothing downstream changes.

Loading the JSON at *run time* could not work: annotations are non-type template parameters consumed during constant evaluation, so the data must exist then. Baking keeps everything in the one constant-evaluated pass — and keeps the `static_assert` below possible.

10.4.3 The staleness guard

That `static_assert` checks every JSON key names a real member of the type. A renamed or misspelt field **fails the build**, naming the offending key, instead of silently dropping its metadata.

WATCH OUT

Side-car edits need a **regeneration**, not just a rebuild — the JSON is baked when `rosetta_gen` runs. Re-run `rosetta_gen` and the generator, not only the binding build.

10.4.4 Attaching the JSON yourself

The manifest's `annotations` field is sugar for specialising a customisation point. You can do it directly, from anywhere — another file, a translation unit, after the class:

```
#include "widget.h"           // the untouched class
#include <rosetta/annotate.h>  // for the customization point

template <>
constexpr std::string_view rosetta::ann_json_source<Widget> = R"({
    "title": { "doc": "The widget title" },
    "count": { "doc": "Visible items", "range": [0, 100] },
    "id":    { "readonly": true },
    "mode":  { "combobox": ["fast", "slow"] }
```

```
}>";
```

Same schema, same merge, same effect on every backend. The one rule is that the specialisation must be visible before `rosetta::walk<Widget>()` runs in a translation unit — which is exactly why the manifest route bakes it into `bindings.h`.

10.4.5 Merge precedence

Inline annotations come first in the merged pack, JSON after. So a reader that keeps the last match (`ann::get_or<A>`) lets a JSON entry *override* an inline one of the same kind, while cumulative readers such as `doc` see both.

A class with no side-car uses the empty primary template and is simply un-annotated — no error, no special-casing.

10.5 Internal annotations

`rosetta::virtual_spec` is **synthesised by the walk**, never written by users: it marks a method's reflection as virtual or overriding so backends can emit overridable bindings (pybind11 and N-API trampolines). You only meet it when writing a custom backend.

10.6 Gotchas

- **Inline annotations make the header C++26-only.** Every translation unit that includes it must build with the fork and `-fannotation-attributes`. Use the side-car when you want off-the-shelf builds.
- **Side-car edits need a regeneration**, not just a rebuild.
- `widget::slider` without a `range` falls back to the default editor — a slider needs bounds.
- `widget::radio` does nothing without a combobox. It restyles the choices; it does not define them.
- `combobox` is capped at 16 choices. Beyond that a drop-down is the wrong widget anyway.
- `label` changes only the visible text. Scripting backends still bind the C++ identifier.

NOTE

The side-car parser is deliberately small: it takes string values literally (no escape unescaping) and parses numbers as plain decimals. Enough for the annotation schema, not a general JSON library.

Part IV

Driving the tool

Chapter 11

The `rosetta_gen` tool

`rosetta_gen` is the single binary that drives the whole pipeline. It reads a manifest and can take you from “a folder of C++ headers” to “compiled bindings for every declared language” in one command — or one step at a time, when you want to inspect what each stage produces.

The tool itself uses **no reflection**, only JSON parsing and text templating, so it builds with an off-the-shelf C++17 compiler. Only the *generator* it emits — and the `rest` and `json` targets — needs the clang-p2996 toolchain.

11.1 Building it

One-time, from the repository root (CMake $\geq$ 3.28; it fetches the single-header nlohmann/json automatically):

```
cmake -G Ninja -S tools/rosetta_gen -B tools/rosetta_gen/build
cmake --build tools/rosetta_gen/build
```

The binary always lands in `<repo>/bin/rosetta_gen`, not in the build tree. If you scaffolded your project with `tools/rosetta_init.py`, its bootstrap `CMakeLists.txt` does this for you and the binary appears at `extern/rosetta/bin/rosetta_gen`.

11.2 The four modes

```
rosetta_gen <manifest.json> [out_dir]      # plain: emit the generator project only
rosetta_gen --build <manifest.json> [opts]  # the whole pipeline in one command
rosetta_gen --clean <manifest.json> [opts]  # remove everything generated
rosetta_gen --init [manifest.json] [src]    # write a starter manifest
```

`rosetta_gen -help` prints the full built-in help; `-build -help` and `-clean -help` list each mode’s options.

```
--init --> manifest.json --> --build -----> bindings/python node wasm ...
(start)   (you edit it)      |                                     ^
                               | = the three manual steps:|
                               | 1. emit gen/           |
                               | 2. cmake gen/, run it --+
                               | 3. cmake / npm / emcmake each backend
--clean <-- removes gen/ (or generated/), the generator binary and bindings/
```

11.3 `-init`: starting a project

Writes a starter manifest (default `./manifest.json`). It **never overwrites** an existing file — remove it first to regenerate.

A blank, commented manifest

```
bin/rosetta_gen --init
```

Emits a fully-commented example exercising every commonly-used field: a `variables` declaration feeding the `cpp26_*` overrides, a multi-entry `user_include`, `user_sources`, `compile_definitions`, `build_type` / `optimization`, a representative spread of targets, one class and one function. The idea is that you delete what you do not need rather than hunting the documentation for what exists.

Pre-filled from a source scan

```
bin/rosetta_gen --init manifest.json ./src
```

An argument naming an existing **directory** is the source tree to scan; the other, if any, is the manifest path. The scan walks for headers (`.h/.hh/.hpp/.hxx`) and sources (`.cpp/.cxx/.cc`) and pre-fills:

- `classes` and `functions`, from the declarations found, with `shared namespace` / `header_dir` defaults factored out when every entry agrees;
- `user_sources` — every source file found, **except** any defining `main()` (it would clash with the binding modules' entry points; reported as a note);
- `rosetta_include`, guessed by walking up from the `rosetta_gen` binary; a `FIXME` comment otherwise.

| WATCH OUT

The scan is a comment- and string-aware *token scan*, not a real C++ parse. Template classes, overloaded free functions and anonymous namespaces are skipped, each with a printed note. An overloaded free function is skipped because the scan has no signature to write, not because it cannot be bound — add it by hand with a **signature** (Section 5.4.1). **Review the result before building.**

VCS directories, build trees and vendored code (`build/`, `extern/`, `third_party/`, `node_modules/`, ...) are not descended into.

| NOTE

For a whole *project* skeleton — a bootstrap `CMakeLists.txt` that fetches `rosetta` into `extern/`, a `.gitignore`, a `README.md` — rather than just the manifest, use `tools/rosetta_init.py`.

11.4 `-build`: the whole pipeline

```
bin/rosetta_gen --build manifest.json
```


runs, in order:

1. **emit** the generator project into `<manifest dir>/gen/`;
2. **build** it with CMake — this step needs the clang-p2996 toolchain, and reflection runs here — and **run** the resulting generator, writing one self-contained project per target into `<manifest dir>/bindings/<lang>/`;
3. **compile every backend** the manifest declares, each with its own build shape.

Backend	Build
python, nanobind, rest, json, lua, imgui, ...	plain cmake configure + build
node	npm install + npm run build
wasm	emcmake cmake + cmake -build
julia	cmake (locates JLCxx by running julia)
csharp	cmake, then dotnet build as a second stage
java	cmake, then mvn package as a second stage
qt, qml	cmake with -DQT_DIR
markdown, html, typescript, openapi, paraview	nothing to compile

A backend whose toolchain is **missing** — no npm, no emcc, no julia — is *skipped with a note*, not an error. Likewise a missing dotnet or mvn downgrades csharp / java to “OK (native only)”. A backend whose build **fails** is reported, and -build exits non-zero *after trying the rest*.

Every run ends with a per-backend summary:

```
== summary
python      OK
node        OK
wasm        SKIPPED (emcc not found -- activate emsdk)
markdown    OK (nothing to compile)
```

11.4.1 Options

```
--gen-dir DIR          generator project dir   (default: <manifest dir>/gen)
--bindings-dir DIR     generated bindings dir  (default: <manifest dir>/bindings)
--clang-p2996-root PATH -DCLANG_P2996_ROOT for every CMake configure
--qt-dir PATH          -DQT_DIR for the qt / qml builds
--only a,b / --skip a,b restrict the backend builds to / exclude these
--cmake-arg ARG        extra argument for every CMake configure (repeatable)
--jobs N, -jN          parallel build jobs (bare -j: one per core)
--fresh               wipe the gen and bindings dirs first
--wheel               also build a Python wheel for the python /
                       nanobind targets
--wheel-dir DIR        collect every wheel in DIR (implies --wheel)
```

11.4.2 Iterating on one or two backends

While developing you rarely want the full matrix. -only and -skip filter **which backends get compiled** — generation still emits them all, and the filters name the lang values:

```
# just python and node, 8-way parallel
bin/rosetta_gen --build manifest.json --only python,node -j8
```

```
# everything except the slow wasm link
bin/rosetta_gen --build manifest.json --skip wasm
```

11.4.3 A custom reflection toolchain

rest is the one target whose generated code still splices reflections, so it re-runs reflection when *it* compiles. Its CMake configure needs to know where the clang-p2996 build lives if it is not at the manifest's default:

```
bin/rosetta_gen --build manifest.json \
  --clang-p2996-root ~/devs/clang-p2996/build --fresh
```

-fresh wipes *gen/* and *bindings/* first — the right move after toolchain or manifest-layout changes, when stale CMake caches would otherwise bite.

11.4.4 Qt and QML

```
bin/rosetta_gen --build manifest.json --only qt,qml --qt-dir ~/Qt/6.8.3/macos
```

Equivalently, set *qt_dir* in the manifest and drop the flag.

11.4.5 Passing flags to every configure

-cmake-arg is repeatable and applies to the generator's configure *and* every backend's, including the emcmake configure for *wasm*:

```
bin/rosetta_gen --build manifest.json \
  --cmake-arg -DCMAKE_BUILD_TYPE=Debug \
  --cmake-arg -DCMAKE_EXPORT_COMPILE_COMMANDS=ON
```

11.4.6 Python wheels

By default *-build* compiles the extension module and stops — it produces an importable *.so* / *.pyd*, not a distributable artifact. *-wheel* adds the packaging step for the two backends that emit a *make_wheel.py*:

```
bin/rosetta_gen --build manifest.json --only python,nanobind --wheel
```

Wheels land in *<bindings>/<lang>/dist*, with the repaired (redistributable) copies in *dist/repaired*. It is opt-in because packaging is a release action, not a development one: it pip-installs *build*, *scikit-build-core* and a repair tool into the environment, and takes noticeably longer than the compile *-build* is otherwise optimised for.

-wheel-dir collects every backend's wheels in one place instead — the shape you want when the next step is uploading them:

```
bin/rosetta_gen --build manifest.json --wheel-dir ./dist
```

```
dist/
  pygeom-0.2.0-cp312-cp312-macosx_15_0_arm64.whl
  nbgeom-0.2.0-cp312-abi3-macosx_15_0_arm64.whl
  repaired/          <- the redistributable copies, same names
```

Sharing one output directory is safe: each `make_wheel.py` repairs only the wheels *it* just produced.

Both flags have manifest counterparts — `wheel` and `wheel_dir` on the `python` / `nanobind` entry of `targets` (Section 7.5) — so a project that always packages says so once. They are per-target, which is what lets a manifest ship the `nanobind` wheel and not the `python` one. The flags still win, and only ever toward packaging *more*: a target spelling `"wheel": false` cannot disarm a `-wheel` you typed, and `-wheel-dir` overrides every target's own directory.

NOTE

The wheel build is *independent* of the compile above it — `scikit-build-core` re-runs the same `CMakeLists.txt` in its own tree, so it neither reuses nor disturbs `<lang>/build/`. The interpreter is probed on `PATH`; if none is found, the backend is recorded OK (no `wheel - no python interpreter found`) rather than failing. A `-wheel` that cannot fire — no such target, or none surviving `-only/-skip` — says so instead of quietly doing nothing.

11.5 Plain mode: one step at a time

```
bin/rosetta_gen manifest.json [out_dir]
```

emits **only** the generator project (default `out_dir`: `<manifest dir>/generated/`) and stops. Useful for inspecting `bindings.h`, hacking on the emitted driver, or wiring the steps into your own build system:

```
# 1. manifest -> generator project
bin/rosetta_gen manifest.json gen

# 2. build the generator (clang-p2996) and run it
cmake -S gen -B gen/build && cmake --build gen/build
./generator bindings           # the binary is dropped NEXT TO the project dir

# 3. build any backend by hand
cmake -S bindings/python -B bindings/python/build
cmake --build bindings/python/build
```

Two footguns the tool warns about:

- a fresh checkout (or a `-clean`) leaves `out_dir/build` unconfigured, so `cmake -build` alone fails — the *configure* step is printed as a **next:** hint;
- if the emitted sources changed but a generator binary from a previous emit still sits next to `out_dir`, running it without rebuilding **silently re-emits the old bindings**. The tool prints a rebuild note when it detects this.

Re-emitting is idempotent: unchanged files are not rewritten, so the next `cmake -build` recompiles nothing.

11.6 `-clean`: back to sources

```
bin/rosetta_gen --clean manifest.json [--gen-dir DIR] [--bindings-dir DIR]
```

Removes everything the pipeline generated for this manifest:

- the generator project directory — the explicit `-gen-dir`, or **both** defaults (`gen/` from `-build` and `generated/` from plain mode);
- the generator binary dropped beside it, including per-config subdirectories of multi-config generators;
- the per-backend folders under `bindings/` **that the manifest declares**. Anything else you parked there survives, and `bindings/` itself is only removed once empty.

Never the manifest — and never a folder whose `CMakeLists.txt` lacks the `Generated by rosetta` stamp the emitters write, so a hand-written folder sharing a name survives:

```
not removing gen (no rosetta_gen marker in its CMakeLists)
```

11.7 Exit status

Mode	0	non-zero
plain	project emitted	manifest missing / invalid, emit error
-build	every <i>attempted</i> backend built (skips do not count)	generator build or run failed, or at least one backend build failed
-clean	clean ran (nothing to remove is fine)	manifest missing / invalid
-init	manifest written	target file already exists, bad arguments

Chapter 12

Overloads and coverage

Two features that belong together, because the second is what makes the first readable: rosetta binds **overload sets**, and every generation writes a machine-readable `coverage.json` saying what bound, what did not, and why.

12.1 Overloads

The reflection walk deduplicates member functions by **(name, signature)**, so of

```
struct Grid {
    double at(int i) const;
    double at(int i, int j) const;
};
```

the whole set arrives in the IR. What happens next depends on the target language, because “two functions with one name” is not something every runtime can express.

Target	Policy	Why
python	all	repeated <code>.def("at", ...)</code> builds one overload set; pybind dispatches on argument types
nanobind	all	same model as pybind11
julia	all	each becomes a Julia method; multiple dispatch picks
qt	all	builds an independent widget row per method — nothing is keyed by name
wasm	first only	a second <code>.function("at", ...)</code> throws <code>BindingError</code> at module init
node	first only	property descriptors are keyed by name; JS has no type dispatch
lua	first only	<code>c["at"] = ...</code> is an assignment — a second one overwrites
csharp / java	first only	the op table is a name-keyed map, and marshalling goes through JSON
typescript	first only	the <code>.d.ts</code> describes the N-API module, so it must not promise what <code>node</code> did not bind
rest	first only	a route is a URL path; two overloads would register the same path
qml	first only	<code>registerInvoker</code> is keyed by name, and QML calls with an untyped <code>QVariantList</code>

“First” always means **first-declared**, which is stable across regenerations.

Every dropped overload appears in `coverage.json` with the reason `overload_not_expressible`, so you can see exactly what a name-keyed target left behind — and reach it with an extension method under a distinct name if you need it.

NOTE

This is also the one thing the runtime object model buys back. Through the `dynamic` projection, those same five name-keyed targets get their overloads *back*, because resolution happens at call time rather than at generation time. See Chapter 14.

12.1.1 C++ name hiding is honoured

```
struct Base { int f(int) const; int f(int, int) const; };
struct Derived : Base { int f(double) const; };
```

`Derived` binds **only** `f(double)`. That is what a C++ caller sees without a `using Base::f`, so binding the hidden base overloads would hand scripts a call the C++ API does not offer. Both hidden declarations are recorded as `hidden_by_derived` drops rather than vanishing.

A diamond-shared base member is still emitted exactly once. Two *different* bases declaring the same name both bind — that call is ambiguous in C++ without qualification, but the binding has no ambiguity to resolve (each entry is spliced from its own reflection), so binding both beats binding neither.

12.2 coverage.json

`generate()` writes it to `<out_dir>/coverage.json` on every run, always. An empty skip list is itself the useful answer, and a file that is always present is one you can diff.

```
{
  "rosetta_coverage": 1,
  "reflection": [                                // never reached ANY backend
    { "class": "Shape",
      "dropped": [
        {"member": "operator<", "signature": "bool(const Shape &) const",
         "reason": "no_identifier"},
        {"member": "debug", "signature": "", "reason": "function_template"}
      ]
    }
  ],
  "targets": [
    { "target": "wasm",
      "bound": 3, "skipped": 1,
      "classes": [
        { "class": "Shape",
          "bound": [ {"member": "area", "signature": "double () const"} ],
          "skipped": [ {"member": "scale", "signature": "void (double, double)",
                       "reason": "overload_not_expressible",
                       "detail": "the target binds methods by name and ..."} ]
        }
      ]
    }
  ]
}
```

12.2.1 The two stages

reflection lists members the walk never handed to any backend. These are backend-independent:

Reason	Meaning
<code>no_identifier</code>	an operator or conversion function — no name to bind to
<code>function_template</code>	a member template: no fixed parameter pack to splice
<code>hidden_by_derived</code>	a base overload hidden by a derived declaration of the name

targets is what each backend's own gates decided:

Reason	Meaning
<code>unmarshalable_signature</code>	no conversion for a type in the signature; detail names it
<code>overload_not_expressible</code>	the target keys methods by name and this one lost the tie
<code>sequence_not_adaptable</code>	a registered sequence with no <code>std::vector</code> boundary adapter here
<code>static_callback_receiver</code>	<code>embind</code> cannot give a static member the receiver a callback adapter needs
<code>extension_has_no_member_pointer</code>	this backend dispatches through a member pointer

Both the **bound** and the **skipped** side are recorded. Counting only skips cannot tell “everything bound” apart from “this class never reached this backend” — and that is exactly the distinction that matters when a whole target silently produces nothing.

Free functions get their own “functions”: `{"bound": [...], "skipped": [...]}` array beside “classes”. On a library whose algorithms are free functions over one data structure, that is most of the API.

12.3 Using it

The point is the *diff*. Check `coverage.json` into your repository next to the manifest, and a method that stops binding becomes a reviewable line in a pull request instead of an `AttributeError` a week later.

To gate CI on it:

```
jq -e '[.targets[].skipped] | add == 0' bindings/coverage.json
```

More usefully: allow known skips and fail on new ones by diffing against a committed baseline. A binding surface that only ever grows deliberately is worth the twenty lines of CI it costs.

| WATCH OUT

One asymmetry to know before comparing numbers across targets: **typescript** and **dynamic** record bound *fields*, while the other language backends record only methods. The top-level bound scalar is therefore not comparable across those two groups — count methods by signature instead.

12.4 Instrumenting a backend of your own

Skips are recorded at the existing gate sites, one line each:

```

for (const auto &m : k.methods) {
    if (!coverage::emit_overload(coverage::overloads::first_only,
                                "my-backend", k, m)) {
        continue; // records the dropped overload
    }
    if (!my_method_ok(m, c)) {
        coverage::note_skip("my-backend", k, m, "unmarshalable_signature",
                            my_skip_reason(m, c)); // a slug plus a sentence
        continue;
    }
    segs.push_back(my_method(k, m));
    coverage::note_bound("my-backend", k, m);
}

```

Free functions use `note_bound_function` / `note_skip_function`, fields `note_bound_field` / `note_skip_field`.

NOTE

Keep the reason function next to the predicate it explains. A reason that drifts out of step with the decision is worse than no reason at all.

12.5 A note for backend authors

`&T::name` names an **overload set**, not a function, so any emitter that spells a member pointer for an overloaded name must cast it to the exact signature. Use the shared helper:

```

const std::string mp = px_member_pointer(k, m); // &T::f, or a static_cast

```

The gate is `GenMethod::is_overloaded` — does the *C++ class* overload the name? — and `not overload_count`, which is how many entries reached the IR. A set whose siblings were all gated out still needs the cast.

Part V

Beyond bindings

Chapter 13

The dynamic object model

Every backend so far spends the reflection result on *framework calls*: `py::class_<Mesh>().def("at", ...)`, an N-API property descriptor, a sol2 usertype. The information is consumed at generation time, and what survives is code for one specific runtime.

The dynamic backend spends it on **data instead**. It writes the IR back out as static tables, so the answer to “what does this class have” is still available at run time — to a caller that was never compiled against your headers.

13.1 What it emits

<code>bindings/dynamic/auto_dynamic.h</code>	declares <code><lib>::register_all()</code> -- idempotent
<code>bindings/dynamic/auto_dynamic.cpp</code>	the tables
<code>bindings/dynamic/inspect.cpp</code>	a registry walker, for free
<code>bindings/dynamic/CMakeLists.txt</code>	a static library, <code><lib>_dynamic</code>

`auto_dynamic.cpp` is the only file that includes your header. It is ordinary C++20: the metadata is data, so nothing here needs reflection to compile.

13.2 Using it

```
#include <rosetta/dynamic.h>    // the runtime -- ordinary C++20
#include "auto_dynamic.h"       // the generated tables (scene.h is NOT included)

scene::register_all();          // publish the tables into the registry, once

for (const rosetta::dyn::MetaClass *k : rosetta::dyn::registry().classes())
    for (std::size_t i = 0; i < k->n_fields; ++i)
        std::cout << k->qualified << "::" << k->fields[i].name << '\n';
        // ... plus .type, .doc, .readonly, .range, .choices
        // and the .get / .set thunks
```

Note what the consumer does *not* do: it never includes `scene.h`, and it names no bound type. That is the property everything else in this chapter follows from.

13.3 What the tables look like

One `TypeDesc` per distinct type, shared by address; one `MetaField` / `MetaMethod` / `MetaCtor` array per class; one `MetaClass` tying them together. Straight from generated output:

```

TypeDesc td_0{.kind = Kind::number, .spelling = "double", .integral = false};
TypeDesc td_6{.kind = Kind::vector, .spelling = "std::vector<double>",
              .element = &td_0};

const MetaField k_scene__Vec3_fields[] = {
    {.name = "x", .type = &td_0, .doc = "X component",
     .annotations = k_scene__Vec3_x_anns, .n_annotations = 2,
     .get = +[](const ObjectRef &self, const ArgList &) {
         return make_any(&td_0, static_cast<scene::Vec3 *>(self.ptr)->x);
     },
     .set = +[](const ObjectRef &self, const ArgList &a) {
         static_cast<scene::Vec3 *>(self.ptr)->x = value_cast<double>(a[0]);
         return Any::none();
     }},
};

```

Three properties of that shape are deliberate:

- **Aggregate initialisers of trivially-constructible data** — `const char *`, pointer-and-count spans, function pointers. No `std::string`, no `std::vector`, no dynamic initialisation. A 500-class library’s metadata lands in `.rodata`, not in a static constructor that runs at load.
- **One captureless lambda per member**, decaying to a plain function pointer: using `Thunk = Any (*)(const ObjectRef &self, const ArgList &args)`. That is the *only* callable shape in the model — a field getter, a field setter, an instance method, a static method and a free function all reduce to it. `self` is the receiver, default-constructed when there is not one.
- **TypeDesc is structural, not an identity**. `std::type_index` answers “are these the same type”; a UI and a marshaller need “what is *inside* it”, so `element` recurses and `vector<vector<Foo>>` is describable.

13.4 Why `register_all()` exists

Two classes can name each other’s types, so `TypeDesc::cls` cannot be filled during static initialisation without an ordering you have no way to guarantee. The generated function adds every class, enum and function to the registry and then calls `link()` to resolve those back-pointers:

```

void register_all(rosetta::dyn::Registry &r) {
    r.add_class(&k_scene__Vec3);
    r.add_class(&k_scene__Mesh);
    r.link();
}

```

It is idempotent, and it is the one line a consumer must run before querying anything.

13.5 Calling by name

`Any` is the value type crossing the boundary, and its **canonical representation** is what makes a single generic wrapper possible: whatever the C++ said, the box holds exactly one of `bool`, `long long`, `double`, `std::string`, `std::vector<Any>` or an `ObjectRef`.

So a caller reads a number with `as_number()` without caring whether the signature said `int`, `size_t` or `uint32_t` — and changing one for the other in your library is not a break for the wrapper.

`resolve()` scores every candidate of a name against the actual arguments on a deliberately small ladder — `exact` / `promote` / `convert` / `none` — rather than reproducing C++’s conversion lattice. The question from a dynamically typed host is “did the user mean this overload”, not “which standard conversion sequence is shorter”.

An equal best score is reported as *ambiguous* rather than silently picked, and a failure names every candidate and why it lost:

```
Mesh::at(std::string) matched no overload of 2:
  at(int) -- argument 1 ("nope") is not a int
  at(int, int) -- takes 2 argument(s), 1 given
```

13.6 Two consequences worth knowing

What cannot be marshalled is described, not deleted

A member whose type no gate claimed keeps its entry, with a null thunk and a stated reason:

```
.skip_reason = "callback: a std::function parameter needs a foreign-callable adapter"
```

So a user interface can grey it out instead of the member silently vanishing. This is the same information `coverage.json` carries, available at run time rather than at review time.

Ownership is explicit

`ObjectRef::owner` is non-null when the reference owns the object. A getter returning `T&` can hand back a *borrowed* handle carrying the parent’s owner — which pins the parent and closes the dangling-sub-object hole that a raw-pointer return has.

13.7 What it buys

Relative to the language backends:

- **Overloads come back.** The metadata keeps the whole overload set and scores it against the actual arguments, so the name-keyed targets’ first-only limit (Chapter 12) does not apply — and a failed call can name every candidate it rejected, with the reason.
- **UI by query, not by codegen.** A property inspector is a walk over `registry(): label` to caption, `doc` to tooltip, `range` to slider bounds, `readonly` to a disabled row, `combobox` and `enumerators` to a drop-down, `button` to an action row. Annotations are enforced once, in the core, instead of per backend.
- **Lifetime is explicit**, as above.

13.8 What it costs

Every access is a linear name lookup plus a scoring pass.

That is the right price for menus, property panels, commands and scripting glue. It is the wrong price for bulk data — pulling a mesh across boxes every coordinate into an `Any`.

NOTE

Cache the resolved `MetaMethod`: the pointer is stable for the process. And keep an expanded binding for hot paths — the two projections are not exclusive, and a library that emits both gets the property sheet from one and the throughput from the other.

13.9 When to reach for it

You want...	Use
Python/JS/Lua callers of your API	a language backend
a property panel over your types	<code>dynamic</code> + a walker, or the <code>qt</code> / <code>qml</code> / <code>imgui</code> backends
a command console, a menu system, an undo stack keyed by member name	<code>dynamic</code>
one scripting wrapper that works for <i>any</i> library you have bound	<code>dynamic</code> + the <code>scriptable</code> preset (Chapter 14)
maximum throughput on a hot call	a language backend

The worked example is `examples/dynamic`: one set of generated metadata driving a terminal interpreter *and* a Qt viewer — 3-D view, property panel, console — neither of which names a single bound type.

Chapter 14


Scripting the metadata itself

The `dynamic` backend makes your library’s metadata available *to C++*. Every consumer of it in the repository is C++: the terminal interpreter, the Qt viewer, the ImGui runtime, the QML reflected object — hand-written walkers, one per toolkit.

The last step is to bind *that API* through rosetta’s own backends, so the walker can be written in Python, Lua or JavaScript instead. This is the Graphite/GOM move: expose the object model itself to the scripting layer, and the user interface becomes a script rather than a backend.

14.1 The trick, in one sentence

Run the pipeline twice, and make the first run’s *output* the second run’s *input library*.



By stage 2 your types are **data**, not C++. So stage 2’s manifest never names one, and the resulting module works for *any* library that has been through stage 1.

14.2 Why a facade rather than binding `dynamic.h`

`rosetta::dyn` cannot be fed to the backends as it stands. Three shapes block it:

In <code>dynamic.h</code>	Why it blocks	In <code>script.h</code>
<code>MetaField *fields;</code> <code>size_t n_fields</code>	pointer-and-count, so vector marshalling never fires	<code>std::vector<FieldInfo></code> <code>fields()</code>
<code>const MetaClass *,</code> <code>const TypeDesc *</code>	raw pointer to an unbound type — the gate skips it	value-semantics handles
<code>using Thunk = Any</code> <code>(*)(...)</code>	function pointers marshal nowhere	ordinary <code>get / set / call</code>

<rosetta/script.h> is that reshaping and nothing more. Each handle wraps **one const pointer** into the tables, so binding it copies no metadata — the `.rodata` guarantee survives.

NOTE

Nothing in the facade is named `Class`, `Type`, `Object` or `Method`, because those collide in the target languages (`java.lang.Class`, JavaScript's `Object`, Lua's `type`). Hence `ClassInfo`, `TypeInfo`, `Instance`, `MethodInfo`.

Errors come back as `Outcome` (an ok flag plus text), never as an exception. Unwinding through a foreign VM's frames is how you get a crash instead of a `TypeError`.

14.3 The one hand-written piece

`Value` is type erasure — one C++ type standing for “whatever the host language just handed us”. Reflection can *describe* it but cannot know that a Python `float`, a Lua number and a JavaScript `Number` should all become the same box. That mapping is per-language knowledge, so it is written per language, in `runtime/script/{python, lua, node, wasm}.h`.

Only half of each caster is per-language, though:

- **Reading** a host value is irreducible — only Python knows what `PyFloat_Check` means.
- **Writing** one is the same shape everywhere, so it lives once in `rosetta::script::visit()`, and each caster supplies a sink of nine small methods (`on_none`, `on_bool`, `on_int`, `on_real`, `on_string`, `on_enum`, `on_list`, `on_object`, `on_unknown`).

No two casters can drift apart on whether an enum crosses as its integer, or on what an unclaimed type renders as.

The casters reach the generated code through `module_init`'s `headers`, which is the only hook that emits an `#include` *after* the framework header and *before* the bound classes — exactly the window a type caster needs. Because that list is global rather than per-target, each file guards its whole body on a macro only its own target defines.

WATCH OUT

One subtlety the type gate forces: `Value` must stay in the bound `classes` list. Bindability is decided at generation time and cannot know a caster will exist, so dropping `Value` would silently skip every member that mentions it — `Outcome::value`, `Instance::set`, `Instance::call`, all of it. The cost is one vestigial `Value` type per module that nothing ever produces.

14.4 What a project has to write

Nothing, beyond where its own files are. The binding surface is fixed, so it ships as manifest data in `include/rosetta/presets/scriptable.json` and a manifest names it:

```
{
  "preset": "scriptable",
  "rosetta_include": "../..../include",
  "user_include": ["../lib", "../lib/bindings/dynamic"],
  "user_sources": ["../lib/bindings/dynamic/auto_dynamic.cpp"],
  "module_init": { "headers": ["auto_dynamic.h"],
    "statements": ["bank::register_all()"] },
  "targets": ["python", "lua", "node"]
}
```

Six keys, every one about *this* project. Read them against the diagram: `user_sources` is stage 1's output, and `module_init` is the `register_all()` that publishes it.

14.5 What it buys

- **One wrapper per language instead of one per class.** Ship the module once; any library that has been through the dynamic backend becomes scriptable with no further code generation.
- **UI by query, not by codegen.** A property sheet is a walk over `fields()`: doc to tooltip, `range` to slider bounds, `combobox` and enumerators to a drop-down, `readonly` to a disabled row. The dispatch that a C++ property panel spends a file on is about ten lines of Python — editable without recompiling anything.
- **Overloads come back** for the name-keyed targets. Lua, Node, Wasm, C# and Java bind the first overload of a name and drop the siblings at generation time; through this module they do not, because `resolve()` scores the whole set at call time and can explain a failure.
- **Annotations are enforced once**, in the core, rather than re-implemented per backend — a `range` violation comes back as the same message in every language.

That third point is worth dwelling on, because it is the sharpest evidence that the IR is genuinely target-neutral. The same reflected overload set survives into both projections; the expressiveness difference between them is attributable entirely to *when* names are resolved, not to information lost on the way.

The set of C++ members a target can express is not a property of the target language. It is a property of when the target resolves names.

14.6 Running the pipeline twice

`examples/plain-binding` runs it **once**: one library, one manifest, four language bindings. That is the base case, and it is what most of this book describes.

`examples/scriptable-model` runs stages 1–2 once per *manifest*, and it has two:

1. **The user library to metadata.** The dynamic backend turns `scene.h` into `auto_dynamic.{h,cpp}` — the same information the other backends spend on framework calls, emitted as plain data instead.
2. **<rosetta/script.h> to bindings.** A second pipeline that binds rosetta's own reflection API, taking stage 1's output as its `user_sources`.

Practically: two manifests, two `rosetta_gen -build` invocations, in that order. Stage 2 must be re-run when stage 1's *output location* changes — but not when your library's *classes* change, which is the whole point.

Chapter 15

Writing a backend

Rosetta ships nineteen backends, but the generator is **open for extension**: you can teach it a new target without editing `rosetta::generate` or any core file. This chapter adds a Lua backend as a worked example — Lua because it is small enough to show whole, and because the real one exists to compare against.

15.1 The mental model: two layers

A backend is two independent pieces that run at two different times, and keeping them straight is most of the work.

Layer	File you write	Runs when	Reflection?
Runtime visitor	<code>rosetta/visitors/lua.h</code>	when the <i>generated</i> binding is compiled	yes
Generation backend	a <code>rosetta::Backend</code> subclass	when the driver runs	no

The split exists because `rosetta::generate<Ts...>` erases the type pack up front into plain strings:

```
struct GenClass { std::string name; std::string header; std::string doc; };
```

So a Backend never sees a C++ type — only {name, header, doc} per class. All the real reflection happens later, in your *visitor*, when the generated `auto_lua.cpp` is compiled.

```
manifest.json -(rosetta_gen)-> gen/ -(build)-> ./generator -(run)-> bindings/lua/
                                     |
                                     | auto_lua.cpp
                                     | CMakeLists.txt
                                     |
                                LuaBackend::emit() (compiling auto_lua.cpp
                                                         runs the walk)
```

15.2 Layer 1: the runtime visitor

This is where binding actually happens. It implements the **walk visitor concept**: three required member templates plus one optional, each receiving a `std::meta::info` of the member and its annotation pack.

```
// include/rosetta/visitors/lua.h
#pragma once
```

```

#include <experimental/meta>
#include <rosetta/walk.h>
#include <sol/sol.hpp>
#include <string>

namespace rosetta {

    template <typename T> struct LuaVisitor {
        sol::usertype<T> &ut;

        // A data member. 'Fld' is the reflected field; 'Anns...' is its
        // annotation pack (rosetta::doc / readonly / range / ...).
        template <std::meta::info Fld, auto... Anns> void field(const char *name) {
            using F = [:std::meta::type_of(Fld):];
            if constexpr (ann::has<readonly>(Anns...)) {
                ut[name] = sol::readonly_property(
                    [] (const T &s) -> F { return s.[:Fld:]; });
            } else {
                ut[name] = sol::property(
                    [] (const T &s) -> F { return s.[:Fld:]; },
                    [](T &s, F v) { s.[:Fld:] = v; });
            }
        }

        template <std::meta::info Fn, auto... Anns> void method_instance(const char *name) {
            ut[name] = &[:Fn:];
        }

        template <std::meta::info Fn, auto... Anns> void method_static(const char *name) {
            ut[name] = &[:Fn:];
        }

        // OPTIONAL: constructors. Omit the method entirely and the walk skips
        // it -- it is detected with 'requires'.
        // template <std::meta::info Ctor, auto... Anns> void constructor() { ... }
    };

    // Entry point the generated code calls, one per class.
    template <typename T> void bind_lua(sol::state &lua, const char *name) {
        sol::usertype<T> ut = lua.new_usertype<T>(name);
        LuaVisitor<T> v{ut};
        walk<T>(v);
    }

} // namespace rosetta

```

The annotation helpers you will use inside a visitor:

- `rosetta::ann::has<readonly>(Anns...)` — is annotation A present?
- `rosetta::ann::get_or<doc>(doc{"", Anns...)` — fetch annotation A, or a fallback.

And the reflection splices: `[:std::meta::type_of(Fld):]` for a field's type, `obj.[:Fld:]` to access it, `&[:Fn:]` for a member-function pointer.

NOTE

If your target cannot marshal some parameter or return type, mirror the `napi_supported / rest_supported` pattern — a `constexpr` predicate plus an `if constexpr` guard in `field` and

`method_*` — so unbindable members are *skipped* instead of failing the build. And record them: see Section 12’s instrumentation recipe, which is one line per gate.

15.3 Layer 2: the generation backend

This scaffolds the per-target project. It implements:

```
struct Backend {
    virtual ~Backend() = default;
    virtual void emit(const GenContext &) const = 0;
};
```

with everything it needs handed over in `GenContext`:

<code>out_dir</code>	root of the generated tree (<code>bindings/</code>)
<code>lib</code>	this target’s module/library name, from the manifest entry
<code>classes</code>	<code>std::vector<GenClass></code> — {name, header, doc} per class
<code>user_include</code>	absolute directory holding the class headers
<code>rosetta_include</code>	absolute path to rosetta’s <code>include/</code>

A minimal backend writes one `.cpp` that includes the visitor and calls `bind_lua<T>` per class, plus a `CMakeLists.txt`:

```
#include <rosetta/generate.h>
#include <filesystem>
#include <fstream>

struct LuaBackend : rosetta::Backend {
    static void write(const std::filesystem::path &p, const std::string &s) {
        std::filesystem::create_directories(p.parent_path());
        std::ofstream(p) << s;
    }

    void emit(const rosetta::GenContext &c) const override {
        std::string includes, binds;
        for (const auto &k : c.classes) {
            includes += "#include \"" + k.header + "\"\n";
            using rosetta::exposed_of;
            using rosetta::qualified_of;
            binds += "    rosetta::bind_lua<" + qualified_of(k) +
                ">(lua, \"" + exposed_of(k) + "\");\n";
        }

        const auto dir = c.out_dir / "lua";
        write(dir / "auto_lua.cpp",
            "// Generated by rosetta -- do not edit by hand.\n" + includes +
            "#include <rosetta/visitors/lua.h>\n"
            "#include <sol/sol.hpp>\n\n"
            "extern \"C\" int luaopen_" + c.lib + "(lua_State *L) {\n"
            "    sol::state_view lua(L);\n" + binds +
            "    return 0;\n}\n");

        write(dir / "CMakeLists.txt", /* ... mirroring bindings/python's ... */);
    }
};
```

WATCH OUT

Two names, never interchangeable. `qualified_of(k)` gives the C++ spelling — `arch::sinv::Data` — which is what anything the generated translation unit compiles must use. `exposed_of(k)` gives the **host-language** name: the manifest's `expose` override, else the identifier. Getting these the wrong way round is how two same-named bound types become ambiguous in the generated source, and how a `rename` silently fails to take effect.

15.4 Registering it

`rosetta::generate` dispatches each target through a runtime registry keyed by the manifest's `lang` string:

```
std::map<std::string, std::shared_ptr<Backend>> &backend_registry();
void register_backend(std::string lang, std::shared_ptr<Backend> backend);
```

Register at static-init time with the `BackendRegistrar` helper, in one self-contained plugin translation unit:

```
// lua_plugin.cpp -- your plugin, owned entirely by you
#include "LuaBackend.h"

static const rosetta::BackendRegistrar
    lua_reg{"lua", std::make_shared<LuaBackend>()};
```

Because the registrar runs before `main`, the "lua" backend is present by the time the driver calls `generate` — **no edit to generate required**. An unregistered target is warned about and skipped.

NOTE

Your backend does *not* belong in `rosetta::backend`. That namespace is where rosetta's own implementations live, not a registration point. Put yours in your own namespace; what you inherit from is `rosetta::Backend`, and what you call is the public API below.

15.5 Wiring it into a project

Point the manifest at your plugin with `plugins`, and add a target. `rosetta_gen` compiles each listed source into the generated driver's `add_executable`, resolving paths relative to the manifest.

```
{
  "user_include": "./geom",
  "rosetta_include": "../include",
  "plugins": ["lua_plugin.cpp"],
  "targets": [
    { "lang": "python", "name": "pygeom" },
    { "lang": "lua", "name": "mygeom" }
  ],
  "classes": [
    { "name": "Model", "header": "Model.h" },
    { "header": "Point.h" }
  ]
}
```

```
}
```

Then the usual flow:

```
../../bin/rosetta_gen manifest.json gen # gen/ + your plugin in its CMakeLists
cmake -S gen -B gen/build && cmake --build gen/build
./generator bindings # emits bindings/python AND bindings/lua
```

15.6 API reference

Generation layer (<rosetta/generate.h>)

```
struct GenClass { std::string name; std::string header; std::string doc; };

struct GenContext {
    std::filesystem::path out_dir;
    std::string lib;
    std::vector<GenClass> classes;
    std::string user_include, rosetta_include;
};

struct Backend {
    virtual void emit(const GenContext &) const = 0;
};

void register_backend(std::string lang, std::shared_ptr<Backend>);

struct BackendRegistrar {
    BackendRegistrar(std::string lang, std::shared_ptr<Backend>);
};

std::map<std::string, std::shared_ptr<Backend>> &backend_registry();
```

IR accessors — how a backend spells a bound type

- `qualified_of(const GenClass&) / (const GenEnum&)` — the **C++** spelling, for anything the generated TU compiles.
- `exposed_of(const GenClass&) / (const GenEnum&)` — the **host-language** name.
- `exposed_object_of(const GenType&, const GenContext&)` — the exposed name of the type an IR entry *references*.
- `names_type(const GenType&, const K&)` — does this IR type name that bound class or enum?

NOTE

These six lived in `rosetta::gen_detail` before 2026-08. That spelling still compiles — the old names are using-declarations onto these — but it was a “detail” namespace documented as an extension point, so they now sit in `rosetta::` with the rest of the backend API.

Runtime layer (<rosetta/walk.h>) — the visitor concept

- `v.template field<Fld, Anns...>(const char *name);`
- `v.template method_instance<Fn, Anns...>(const char *name);`

- `v.template method_static<Fn, Anns...>(const char *name);`
- `v.template constructor<Ctor, Anns...>();` — **optional**, called only if defined.
- `rosetta::ann::has<A>(Anns...) / rosetta::ann::get_or<A>(fallback, Anns...)`

15.7 Where to copy from

The built-in backends are a declaration header in `include/rosetta/backends/` with its implementation alongside in `include/rosetta/backends/inline/`. They live in `rosetta::backend` — `backend::Python`, `backend::Node` — so the directory, the file and the namespace each say “backend” once and the class says only which one. They are the reference for the CMake flags and the clang-p2996 link line.

The reflection-free header-only runtimes those backends emit `#include` lines for — the N-API marshaller, the C#/Java C-ABI shims, the Qt and ImGui widget kits — live under `include/rosetta/runtime/`. They are compiled by the *generated* project, never by the generator, which is why they are a directory of their own.

The existing visitors in `include/rosetta/visitors/` are the templates to copy for layer 1.

Part VI

Worked examples

Chapter 16

Worked examples

The repository carries eighteen runnable examples. This chapter is a guide to which one answers which question, followed by three walked through in enough detail to copy from.

16.1 The map

Path	What it shows
<code>examples/manifest</code>	Manifest-driven generation for <code>Person</code> , with no class modification. The smallest complete case.
<code>examples/annotate-manifest</code>	Out-of-line annotations from an external JSON file, wired by the manifest's <code>annotations</code> field.
<code>examples/geom-lib</code>	Manifest-driven bindings for a small geometry library — nested types, vectors.
<code>examples/geom-expanded</code>	Ten backends (<code>python</code> , <code>nanobind</code> , <code>node</code> , <code>wasm</code> , <code>qt</code> , <code>qml</code> , <code>csharp</code> , <code>java</code> , <code>lua</code> , <code>julia</code>) built with off-the-shelf toolchains, plus out-of-line annotations. The proof of Section 1.6.
<code>examples/dynamic</code>	The <code>dynamic</code> backend end to end: one set of generated metadata driving a terminal interpreter <i>and</i> a Qt viewer (3-D view, property panel, console), neither naming a bound type.
<code>examples/scriptable-model</code>	The two-stage pipeline of Chapter 14.
<code>examples/plain-binding</code>	One library, one manifest, four language bindings — the base case, run once.
<code>examples/trampoline-python</code>	Overriding C++ virtuals from Python: generated <code>pybind11</code> trampolines from <code>virtual_spec</code> .
<code>examples/trampoline-node</code>	The same from JavaScript, with N-API trampolines.
<code>examples/dynamic-lib</code>	Linking the bindings against a pre-built library — the <code>user_lib</code> case.
<code>examples/moc</code>	The moc-less signals/slots/properties demo built on <code>mini_moc.h</code> .
<code>examples/docgen</code>	A reflection-driven Markdown / HTML reference generator.

Path	What it shows
<code>examples/paraview</code>	ParaView plugin property-panel XML from an annotated <code>vtkThreshold</code> spec — exercises every annotation feature.
<code>examples/qt</code> , <code>examples/qml</code> , <code>examples/imgui</code>	The three inspector backends, each building a form from a reflected struct. <code>imgui</code> uses an <code>Algo.ann.json</code> side-car and builds with an off-the-shelf C++20 compiler, dependencies auto-fetched.
<code>examples/bindings/*</code>	Hand-written reference backends — <code>pybind11</code> , N-API, <code>jlccxx</code> , REST, <code>embind</code> . What the generator emits, written by hand, for comparison.

16.2 The smallest case

`examples/manifest` is the one to read first. An unmodified `person.h`, a manifest of ten lines, and the pipeline. It is Chapter 3 as a runnable directory.

16.3 A geometry library

`examples/geom-lib` is the first realistic shape: several classes, one returning a `std::vector` of another, an enum, and a free function.

The two things to look at:

- **Nested user types cross without being mentioned.** A `Surface` method returning `std::vector<Triangle>` works because `Triangle` is in `classes` — not because anything says how to convert it.
- **One module per target, not one per class.** The generated `pygeom` carries `Model`, `Point`, `Surface`, `Triangle` and `Kind` together.

`examples/geom-expanded` takes the same library and builds ten backends with stock toolchains — no fork anywhere in the second pass. If you want to convince yourself (or a colleague) that the reflection-free claim is real, that is the directory to run.

16.4 The dynamic model, end to end

`examples/dynamic` is the most interesting example in the repository and the one that takes longest to appreciate.

A small scene library — `Vec3`, `Mesh`, and a `Relaxer` solver — goes through the `dynamic` backend once. Two entirely separate consumers then drive it:

- a **terminal interpreter** that lists types, gets and sets fields, and calls methods by name;
- a **Qt viewer** with a 3-D view, a property panel and a console.

Neither names a single bound type. The property panel is a walk over `fields(): doc` becomes a tooltip, `range` becomes slider bounds, `combobox` becomes a drop-down.

What adding a class costs

The example’s own history is the argument. Adding `scene::Relaxer` — tangential Laplacian relaxation, a genuine algorithm with tuning parameters and one method that runs for a while — touched three files, *all of them the library’s own*:

1. `scene.h`, to add the class;
2. `Mesh.ann.json`, for one new accessor’s annotation;
3. one line of `manifest.json`.

No consumer changed. Not the Qt viewer, not the interpreter, not any of the five scriptable drivers — yet all of them gained it: the solver appeared in the type list, its parameters appeared in the property panel with their declared ranges, and its method appeared as a button.

That is what “UI by query rather than by codegen” buys, stated as a diff.

16.5 Rosetta in the wild

Real libraries bound with rosetta, each from a single manifest and with no hand-written wrappers:

Project	Bound library	Targets
<code>pmp-rosetta</code>	PMP — polygon mesh processing (remeshing, smoothing, subdivision, decimation)	Python, Node.js, WebAssembly, TypeScript
<code>geogram-rosetta</code>	geogram — geometry processing (reconstruction, remeshing, parameterisation, booleans)	Python, Node.js, WebAssembly, TypeScript, Lua
<code>arch-rosetta</code>	Arch — 3-D boundary-element geomechanics	Python, Node.js, WebAssembly, TypeScript
<code>cassini-rosetta</code>	Cassini — FEM geomechanical restoration, through a single high-level facade	Python, Node.js, WebAssembly, TypeScript
<code>triax-rosetta</code>	Triax — discrete-element triaxial compression	Python, Node.js, WebAssembly

The geogram manifest is the best worked reference for a difficult library: it uses `sequences` for `GEO::vector`, `extensions` to reach geometry behind raw-pointer accessors, `signature` to pick one overload of `mesh_union`, `module_init` for `GEO::initialize()`, `generated_headers` for the configured `version.h`, and `link_options` for the wasm NODEFS mount — in other words, nearly every field in Part II, each for a reason you can see.

16.6 Extending a generated binding by hand

Everything under `bindings/` is regenerated output and must never be edited. When the generated binding misses something you need — a helper the walker skips, a typed-array export for a renderer — add it in a **separate hand-written file** compiled *alongside* the generated source, from a small build of your own that lives outside `bindings/`.

The binding frameworks accept several registration blocks per module, so your file contributes a second one. Nothing generated is touched, and regenerating never clobbers your extensions.

```
// viz_helpers.cpp -- compiled next to the generated auto_emscripten.cpp
emscripten::val mesh_positions(const pmp::SurfaceMesh &mesh); // Float32Array
emscripten::val mesh_triangles(const pmp::SurfaceMesh &mesh); // Uint32Array
void triangulate_mesh(pmp::SurfaceMesh &m) { pmp::triangulate(m); }

EMSCRIPTEN_BINDINGS(pmp_viz) { // a second block; the generated one keeps its own
    emscripten::function("mesh_positions", &mesh_positions);
    emscripten::function("mesh_triangles", &mesh_triangles);
    emscripten::function("triangulate_mesh", &triangulate_mesh);
}
```

WATCH OUT

Embind is the friendliest here because it accepts any number of `EMSCRIPTEN_BINDINGS` blocks in one module. Backends with a single module entry point — pybind11's `PYBIND11_MODULE`, N-API's `NODE_API_MODULE` — cannot add a second block to the *same* module. There, ship your extras as a small companion module built the same way (it sees the same user headers and links the same library) and import both.

Before reaching for this, check whether `extensions` (Section 5.5) covers your case: a free function taking `Cl&` needs no build of your own and reaches every backend at once.

Part VII

Reference

Appendix A

Manifest field reference

Every top-level field, alphabetically. • marks a required field.

Field	Req.	Default	Meaning
build_type		Release	Default <code>CMAKE_BUILD_TYPE</code> baked into every compiled backend's <code>CMakeLists.txt</code> . §7.3
classes	•*	—	The classes, structs and enums to bind. *Optional when a <code>preset</code> supplies them. Ch. 5
compile_definitions		[]	Preprocessor definitions (" <code>NAME</code> " or " <code>NAME=VALUE</code> ") applied to the generator driver <i>and</i> every compiled binding target. Ch. 7
cpp26_cc		<code>\${cpp26_root}/bin/clang</code>	C compiler of the reflection toolchain.
cpp26_cxx		<code>\${cpp26_root}/bin/clang++</code>	C++ compiler of the reflection toolchain.
cpp26_lib		<code>\${cpp26_root}/lib</code>	Directory holding the fork's <code>libc++ / libc++abi</code> .
cpp26_root		a built-in path	Root of the C++26 / P2996 toolchain. Moves the other three together. §4.5
functions		[]	Free (non-member) functions to bind. §5.4
generated_headers		[]	Headers the bound library's own build system would have produced, written into the generated tree and put first on the include path. §7.4
generator_name		manifest's directory name	Basename of the emitted reflection driver; last-resort default for <code>module_name</code> .
header_dir		—	Directory fragment prepended to every entry header. §4.7
interop		[]	Foreign libraries whose types the target's binding framework marshals itself (<code>["eigen"]</code>). §8.3

Field	Req.	Default	Meaning
<code>matrices</code>		<code>[]</code>	Foreign 2-D containers that marshal as an array of row arrays. §8.2
<code>module_init</code>		—	Statements the generated module runs when it <i>loads</i> , plus the headers declaring them. §5.7
<code>module_name</code>		a preset's, else <code>generator_name</code>	Default binding module name, used by any bare-string target.
<code>namespace</code>		—	Default C++ namespace for entry names carrying no <code>::</code> of their own. §4.7
<code>optimization</code>		—	Explicit optimisation flag applied after the build type's own. §7.3
<code>out_dir</code>		—	Where every target's <i>built artifact</i> is copied after each build. A per-target <code>out_dir</code> overrides it. §6.4
<code>out_params</code>		<code>{}</code>	Which parameters a method returns <i>through a reference</i> , keyed <code>"Class::method"</code> . Never inferred. §5.6
<code>plugins</code>		<code>[]</code>	Extra <code>.cpp</code> sources compiled into the generator driver — a custom backend. Ch. 15
<code>preset</code>		—	A named JSON fragment supplying a fixed binding surface. Ch. 9
<code>rosetta_include</code>	•	—	Path to rosetta's <code>include/</code> directory.
<code>sequences</code>		<code>[]</code>	Foreign sequence containers that marshal like <code>std::vector<T></code> . §8.1
<code>targets</code>	•	—	The language backends to emit. Ch. 6
<code>user_include</code>	•	—	Directory holding your class headers, or an array of directories.
<code>user_lib</code>		—	Pre-built external libraries to link the bindings against; one object or an array. Ch. 7
<code>user_sources</code>		<code>[]</code>	User <code>.cpp</code> / <code>.c</code> files compiled into every binding target. Globs allowed. Ch. 7
<code>variables</code>		<code>[]</code>	Manifest-local variables, substituted as <code>\$NAME</code> / <code>\${NAME}</code> into every string. §4.4
<code>version</code>		<code>0.1.0</code>	Distribution version stamped into the packaging artifacts (PEP 440). §7.5

Keys beginning with `//` are comments and are ignored everywhere, including inside groups and entries.

A.1 Class entry fields

Field	Req.	Meaning
header	•	Filename emitted into <code>#include</code> , resolved against <code>user_include</code> . May be a glob.
name		C++ type name; defaults to the header's stem. May be qualified.
expose		The binding name. Must be a plain identifier.
annotations		Path to an out-of-line annotation side-car.
doc		Description for the document backends.
extensions		Free functions exposed as instance methods.
final		No trampoline even with public virtuals.
exclude		Glob-entry only: patterns whose matches are dropped.
entries		Makes this a <i>group</i> rather than an entry. §4.7.1

A.2 Function entry fields

Field	Req.	Meaning
name	•	Function name, optionally qualified.
header	•	Header declaring it.
doc		Description.
expose		The binding name.
signature		C++ function type of the one overload to bind. Required when <code>name</code> is overloaded, meaningless otherwise.

A.3 Target entry fields

Field	Req.	Meaning
lang	•	The backend. See Appendix C.
name		Module / library name; defaults to <code>module_name</code> .
link_options		Extra linker flags for this target only.
out_dir		Where this target's artifact is copied.
python		Interpreter the binding is built for.
requires_python		Minimum Python version; feeds both CMake and <code>pyproject.toml</code> .
napi_version		<code>NAPI_VERSION=</code> for a Node addon.
node_engine		<code>engines.node</code> in <code>package.json</code> .
wheel		Build a wheel for this target on every <code>-build.python</code> / <code>nanobind</code> only — an error elsewhere. §7.5
wheel_dir		Where this target's wheels go instead of its own <code>dist/</code> . Implies <code>wheel</code> .

A bare string in targets is shorthand for `{"lang": "..."}.`

A.4 `user_lib` entry fields

name	•	Library base name — the <code>space</code> in <code>libspace.dylib</code> .
dir	•	Directory holding it; used for <code>-L</code> and <code>rpath</code> .
link		"shared" (default), "static", or "dynamic".

Appendix B

Annotation reference

B.1 The kinds

Kind	Applies to	Inline / side-car	Meaning
doc	fields, methods, classes*	<code>= rosetta::doc{"..."}</code> <code>"doc": "..."</code>	Description text.
range	numeric fields	<code>= rosetta::range{lo, hi}</code> <code>"range": [lo, hi]</code>	Value bounds; validating setter, clamping slider, documented constraint.
readonly	fields	<code>= rosetta::readonly</code> <code>"readonly": true</code>	Getter-only property.
combobox	string fields	<code>= rosetta::combobox>{"a","b"}</code> <code>"combobox": ["a","b"]</code>	Allowed choices, maximum 16.
label	fields, methods	<code>= rosetta::label{"..."}</code> <code>"label": "..."</code>	Display-name override in UI backends only.
button	methods	<code>= rosetta::button{"Run"}</code> <code>"button": "Run"</code>	Call-button caption.
widget::*	fields	<code>= rosetta::widget::slider</code> <code>"widget": "slider"</code>	Editor hint; see below.

* Class-level and free-function doc come from the manifest's doc key, since free functions carry no in-source annotations.

B.2 Widget hints

Hint	Field type	Renders as
spin	number	spin box / typed input
slider	number + range	slider bounded by the range
textfield	number	validated free-text input
checkbox	number (0/1), bool	checkbox
color	string "#rrggbb"	colour picker; writes back hex
multiline	string	multi-line text area
radio	string + combobox	radio-button group
file	string path	entry + platform file dialog

B.3 Consumption by backend family

Family	Honours
Scripting — python, nanobind, node, wasm, lua, julia, csharp, java	doc (docstrings where the framework has them), range (validating setters), readonly (getter-only properties). The UI kinds are ignored.
UI inspectors — qt, qml, imgui	Everything.
Documents — markdown, html, openapi, typescript	doc and the constraints (range , readonly , combobox) rendered into the reference text or schema.

An annotation a backend has no use for is ignored, never an error.

B.4 Inline versus side-car

	Inline	Side-car
Header stays plain C++	no	yes
Generated binding builds on an off-the-shelf compiler	no	yes
Works on third-party headers you cannot edit	no	yes
Expressiveness	full set	full set
Edit requires	regeneration	regeneration

Wiring a side-car:

```
"classes": [
  { "name": "Algo", "header": "Algo.h", "annotations": "Algo.ann.json" }
]
```

B.5 Side-car schema

Keys are member identifiers — fields or methods. Each value is an object carrying any subset of:

Key	JSON type	Example
<code>doc</code>	string	"Visible items"
<code>range</code>	2-element array	[0, 100], [1e-10, 1e-6]
<code>readonly</code>	boolean	true
<code>combobox</code>	string array	["fast", "slow"]
<code>label</code>	string	"Max iterations"
<code>button</code>	string	"Run"
<code>widget</code>	string	"slider"

A key naming no member of the type **fails the build**, naming the offending key. That is deliberate: a renamed field must not silently lose its metadata.

B.6 Merge rules

- No side-car: the class is un-annotated. Not an error.
- Inline annotations present: the JSON entries are *concatenated* after them.
- Precedence: a reader that keeps the last match lets a JSON entry override an inline one of the same kind; cumulative readers such as `doc` see both.

B.7 Internal

`rosetta::virtual_spec` is synthesised by the walk and never written by users. It marks a method's reflection as virtual or overriding so backends can emit trampolines.

Appendix C

Target reference

C.1 Every lang value

lang	Output	Builds with	Overloads
python	pybind11 extension module	off-the-shelf compiler	all
nanobind	nanobind extension module	off-the-shelf compiler	all
node	N-API native addon	off-the-shelf compiler	first only
wasm	Emscripten / embind module	off-the-shelf emsdk	first only
lua	sol2 module, require -able	off-the-shelf compiler	first only
julia	CxxWrap.jl / jlcpp module	compiler + CxxWrap.jl	all
csharp	C-ABI lib + P/Invoke wrappers	compiler + .NET	first only
java	C-ABI + FFM wrappers	compiler + JDK	first only
qt	Qt Widgets inspector	compiler + Qt 6	all
qml	QML / QtQuick inspector	compiler + Qt 6	first only
imgui	Dear ImGui inspector app	off-the-shelf compiler	—
rest	cpp-http lib server + client	C++26 toolchain	first only
openapi	OpenAPI 3.1 specification	text output	—
json	nlohmann (de)serialisation	C++26 toolchain	—
typescript	ambient .d.ts declarations	text output	first only
markdown	API reference document	text output	—
html	styled API reference page	text output	—
paraview	Server Manager plugin XML	text output	—
dynamic	the IR as runtime data	off-the-shelf compiler	all, at call time

Each name also accepts a deprecated `-expanded` spelling resolving to the same backend.

C.2 Feature support

Which manifest features reach which backend. A blank cell means the backend ignores the field rather than failing.

lang	module_init	out_dir	extensions	sequences	out_params
python	yes	yes	yes	yes	yes
nanobind	yes	yes	yes	yes	yes
node	yes	yes	yes	yes	yes
wasm	yes	yes	yes	yes	yes
lua	yes	yes		yes	yes
julia	yes	yes			
csharp					
java					
qt					
qml					
imgui					
rest					
typescript			yes	yes	yes
markdown			yes		
html			yes		

C.3 Packaging

Only two backends emit a `make_wheel.py`, so only those two accept the per-target `wheel / wheel_dir` (Section 7.5); the field is an error on any other `lang`.

python	One wheel per CPython version — pybind11 has no stable-ABI mode. Covering several means running <code>make_wheel.py</code> once per interpreter.
nanobind	One abi3 wheel covering every CPython 3.12+; below 3.12, tagged for the exact interpreter that built it.
every other	No packaging step.

C.4 Interop and foreign types

lang	interop: ["eigen"]	
python	native numpy, through <code><pybind11/eigen.h></code>	
nanobind	native numpy, through <code><nanobind/eigen/dense.h></code> (dense only)	
every other	skips the member — register the concrete type under <code>sequences / matrices</code> to give them a flat array	

lang	std::filesystem::path	std::shared_ptr<T> returned
python	native (<code>os.PathLike</code> in, <code>pathlib.Path</code> out)	<code>py::class_<T, shared_ptr<T>></code>
nanobind	native	native
wasm	string, via a copy adapter	<code>.smart_ptr<...></code> on the class
lua	string, via a copy adapter	native (sol2)
node	string, via a copy adapter	the JS object adopts the pointer; untrampolinee pointees only
typescript	string	the pointee type

A `shared_ptr<T>` *parameter* is skipped everywhere — take `const T&`.

C.5 Build shapes for `-build`

Backend	Build
python, nanobind, rest, json, lua, imgui	plain <code>cmake</code> configure + build
node	<code>npm install</code> + <code>npm run build</code>
wasm	<code>emcmake cmake</code> + <code>cmake -build</code>
julia	<code>cmake</code> (locates JLCxx by running <code>julia</code>)
csharp	<code>cmake</code> , then <code>dotnet build</code>
java	<code>cmake</code> , then <code>mvn package</code>
qt, qml	<code>cmake</code> with <code>-DQT_DIR</code>
markdown, html, typescript, openapi, paraview	nothing to compile

A missing toolchain is a *skip with a note*, not a failure. A missing `dotnet` or `mvn` downgrades its backend to “OK (native only)”.

C.6 External requirements

python	pybind11 (<code>pip install pybind11</code>)
nanobind	nanobind; wheels are <code>abi3</code> on Python 3.12+
node	<code>npm</code> , and <code>cmake-js</code>
wasm	an activated <code>emsdk</code>
lua	Lua 5.1–5.4 or LuaJIT — not 5.5, which <code>sol2</code> does not support yet. <code>sol2</code> itself is fetched at configure time
julia	<code>CxxWrap.jl</code> ; the build runs <code>julia</code> to locate JLCxx
csharp	a .NET SDK
java	a JDK with the FFM API
qt, qml	Qt 6 — but no <code>moc</code> on the generated code
imgui	nothing; GLFW and OpenGL3 are auto-fetched
rest	<code>cpp-httplib</code> and <code>nlohmann/json</code> , fetched on first configure, plus the C++26 toolchain

Appendix D

Troubleshooting

D.1 A method is missing from the binding

Read `bindings/coverage.json` first. It was written before anything was compiled, and it names every member that did not reach a backend, with a reason. The likely entries:

Reason	What to do
<code>overload_not_expressible</code>	The target keys methods by name and this overload lost the tie. Wrap it in an <code>extensions</code> entry under a distinct name, or reach it through the dynamic model. §5.5
<code>unmarshalable_signature</code>	<code>detail</code> names the offending type. If it is a foreign container or an Eigen type, Chapter 8 is the fix. If it is a <code>std::function</code> or a raw pointer, write an <code>extensions</code> wrapper with a marshallable signature.
<code>no_identifier</code>	An operator or conversion function. There is no name to bind to — expose it as a named extension method.
<code>function_template</code>	A member template has no fixed parameter pack to splice. Add an explicit non-template wrapper.
<code>hidden_by_derived</code>	C++ name hiding, honoured deliberately. Add <code>using Base::f;</code> to the derived class if you want the base overloads visible.
<code>sequence_not_adaptable</code>	A registered sequence with no <code>std::vector</code> boundary adapter on this backend.
<code>extension_has_no_member_pointer</code>	This backend dispatches through member pointers and cannot take extensions. Expected on <code>qt</code> , <code>qml</code> , <code>csharp</code> , <code>java</code> .

If the member is not in `coverage.json` at all, it was not public, or its class is not in `classes`.

D.2 The generated driver does not compile

`no member named 'vec3' in namespace 'geom'`

A glob entry derived a class name from a file stem that names no such class. See the table in Section 5.2: list the class explicitly, or add the header to `exclude`.

An error inside a signature you wrote

The signature is spliced verbatim. Spell the types the way the header does, fully qualified. `rosetta_gen` only checks its shape — the types are the compiler’s business.

An error inside a `module_init` statement

Same cause: those are spliced verbatim into the module entry point. Spell types qualified, and remember there are no arguments available at that point.

The bound header sees a different class than the binding will

A configuration macro is missing from `compile_definitions`. Those are emitted for the *generator driver* as well as the binding targets precisely so the two cannot disagree — if you passed the macro some other way, move it into the manifest.

D.3 A side-car annotation had no effect

- Did you **re-run `rosetta_gen`**? The JSON is baked into `bindings.h` at generation time. A rebuild of the binding alone will not pick it up.
- A key naming no member *fails the build*, so a silent no-effect usually means the annotation reached a backend that ignores it — `label` and `widget` are UI-only. See Appendix B.
- `widget::slider` without a `range` falls back to the default editor. `widget::radio` without a `combobox` does nothing.

D.4 The binding is slow

- Check the build type. Omitting `build_type` now defaults to `Release`, but a `-DCMAKE_BUILD_TYPE=Debug` on the command line wins. An unoptimised `pybind11` call costs roughly eight times an optimised one.
- If you are going through the dynamic model, every access is a linear name lookup plus a scoring pass. Cache the resolved `MetaMethod`, and keep an expanded binding for hot paths. §13
- A foreign container crosses *by copy* through a `std::vector` boundary. For bulk data on the Python family, prefer `interop` — the `numpy` caster avoids the copy where the layout allows.

D.5 The module will not load

- **Missing shared library.** Every `user_lib` dir lands in `BUILD_RPATH` and `INSTALL_RPATH`, so this usually means a dependency was not listed. Add it — and remember the order is the link order.
- **A wasm target linking a native library.** Every `user_lib` entry is taken as `lib<name>.a` built with the *same* emsdk. A native `.so` cannot enter a wasm module.
- **Something the library needs was never initialised.** That is what `module_init` is for. §5.7

D.6 The pipeline produced stale output

- **A generator binary from a previous emit.** If the emitted sources changed but the old generator still sits next to the output directory, running it re-emits the old bindings.

The tool prints a rebuild note — believe it.

- **A `user_sources` glob.** Globs expand when `rosetta_gen` runs. A new `.cpp` does not appear by itself.
- **Stale CMake caches** after a toolchain or layout change. Use `-fresh`.

D.7 `-build` skipped a backend

That is by design: a missing toolchain is a skip with a note, not an error. The summary line says which tool was not found.

```
wasm          SKIPPED (emcc not found -- activate emsdk)
```

A backend that *failed* is reported differently, and `-build` exits non-zero after trying the rest.

D.8 The wrong Python was used

`-cmake-arg -DPython_EXECUTABLE=...` **does not work** — the generated CMake sets that variable with `CACHE ... FORCE` after probing. Pin the interpreter on the target instead:

```
{ "lang": "nanobind", "python": "3.11", "requires_python": ">=3.10" }
```

This also fixes a wheel tagged for the wrong version, since the wheel is built by whichever interpreter runs `make_wheel.py`.

D.9 `-clean` did not remove something

By design in three cases: the manifest is never removed; a folder whose `CMakeLists.txt` lacks the `Generated by rosetta` stamp is never touched; and only the per-backend folders *the manifest declares* are removed, so anything else parked under `bindings/` survives.

```
not removing gen (no rosetta_gen marker in its CMakeLists)
```

D.10 A variable in the manifest did not expand

Only a name declared in `variables` is substituted. `$ENV{HOME}` and any undeclared `${...}` are CMake's and pass through untouched — and a *misspelt* name is indistinguishable from those, so it passes through silently and surfaces as a path that does not exist.

`${NAME}` rather than `$NAME` is worth preferring: `$Px` is one name, `Px`.

D.11 Reporting something

The two artifacts that make a report actionable are `bindings/coverage.json` and the manifest. Between them they say what rosetta thought your API was and what it decided to do about it, which is almost always where the disagreement is.